

Student Performance Predictions

```
student_por = read.csv('student+performance/student-por.csv', sep = ';', stringsAsFactors = TRUE)
student_math = read.csv('student+performance/student-mat.csv', sep = ';', stringsAsFactors = TRUE)
all_data = rbind(student_por, student_math)
summary(all_data)

##   school    sex      age      address famsize Pstatus     Medu
## GP:772   F:591   Min.   :15.00   R:285   GT3:738   A:121   Min.   :0.000
## MS:272   M:453   1st Qu.:16.00   U:759   LE3:306   T:923   1st Qu.:2.000
##                               Median :17.00
##                               Mean   :16.73
##                               3rd Qu.:18.00
##                               Max.   :22.00
## 
##   Fedu      Mjob      Fjob      reason      guardian
## Min.   :0.000  at_home :194  at_home : 62  course   :430  father:243
## 1st Qu.:1.000  health   : 82  health   : 41  home     :258  mother:728
## Median :2.000  other    :399  other    :584  other    :108  other   : 73
## Mean   :2.388  services:239  services:292  reputation:248
## 3rd Qu.:3.000  teacher  :130  teacher  : 65
## Max.   :4.000
## 
##   traveltime   studytime   failures   schoolsup  famsup      paid
## Min.   :1.000  Min.   :1.00   Min.   :0.0000  no :925  no :404  no :824
## 1st Qu.:1.000  1st Qu.:1.00  1st Qu.:0.0000 yes:119  yes:640  yes:220
## Median :1.000  Median :2.00  Median :0.0000
## Mean   :1.523  Mean   :1.97  Mean   :0.2644
## 3rd Qu.:2.000  3rd Qu.:2.00  3rd Qu.:0.0000
## Max.   :4.000  Max.   :4.00  Max.   :3.0000
## 
##   activities nursery higher internet romantic   famrel
## no   :528    no  :209   no  : 89  no  :217  no  :673  Min.   :1.000
## yes:516    yes:835   yes:955 yes:827  yes:371  1st Qu.:4.000
##                               Median :4.000
##                               Mean   :3.936
##                               3rd Qu.:5.000
##                               Max.   :5.000
## 
##   freetime      goout      Dalc      Walc
## Min.   :1.000  Min.   :1.000  Min.   :1.000  Min.   :1.000
## 1st Qu.:3.000  1st Qu.:2.000  1st Qu.:1.000  1st Qu.:1.000
## Median :3.000  Median :3.000  Median :1.000  Median :2.000
## Mean   :3.201  Mean   :3.156  Mean   :1.494  Mean   :2.284
## 3rd Qu.:4.000  3rd Qu.:4.000  3rd Qu.:2.000  3rd Qu.:3.000
## Max.   :5.000  Max.   :5.000  Max.   :5.000  Max.   :5.000
## 
##   health      absences      G1       G2
## Min.   :1.000  Min.   : 0.000  Min.   : 0.00  Min.   : 0.00
## 1st Qu.:3.000  1st Qu.: 0.000  1st Qu.: 9.00  1st Qu.: 9.00
## Median :4.000  Median : 2.000  Median :11.00  Median :11.00
## Mean   :3.543  Mean   : 4.435  Mean   :11.21  Mean   :11.25
## 3rd Qu.:5.000  3rd Qu.: 6.000  3rd Qu.:13.00  3rd Qu.:13.00
## Max.   :5.000  Max.   :75.000  Max.   :19.00  Max.   :19.00
```

```

##      G3
##  Min.   : 0.00
##  1st Qu.:10.00
##  Median :11.00
##  Mean   :11.34
##  3rd Qu.:14.00
##  Max.   :20.00

head(all_data)

##   school sex age address famsize Pstatus Medu Fedu      Mjob      Fjob reason
## 1     GP   F   18       U    GT3      A     4     4 at_home teacher course
## 2     GP   F   17       U    GT3      T     1     1 at_home other  course
## 3     GP   F   15       U    LE3      T     1     1 at_home other  other
## 4     GP   F   15       U    GT3      T     4     2 health services home
## 5     GP   F   16       U    GT3      T     3     3 other   other  home
## 6     GP   M   16       U    LE3      T     4     3 services other  reputation
##   guardian traveltIME studytime failures schoolsup famsup paid activities
## 1   mother          2        2     0      yes    no   no    no
## 2   father          1        2     0      no     yes  no    no
## 3   mother          1        2     0      yes    no   no    no
## 4   mother          1        3     0      no     yes  no    yes
## 5   father          1        2     0      no     yes  no    no
## 6   mother          1        2     0      no     yes  no    yes
##   nursery higher internet romantic famrel freetime goout Dalc Walc health
## 1   yes    yes    no    no    4     3     4    1    1    3
## 2   no     yes    yes   no    5     3     3    1    1    3
## 3   yes    yes    yes   no    4     3     2    2    3    3
## 4   yes    yes    yes   yes   3     2     2    1    1    5
## 5   yes    yes    no    no    4     3     2    1    2    5
## 6   yes    yes    yes   no    5     4     2    1    2    5
##   absences G1 G2 G3
## 1        4  0 11 11
## 2        2  9 11 11
## 3        6 12 13 12
## 4        0 14 14 14
## 5        0 11 13 13
## 6        6 12 12 13

write.csv(all_data, 'student+performance/merged_data.csv', row.names = FALSE)

```

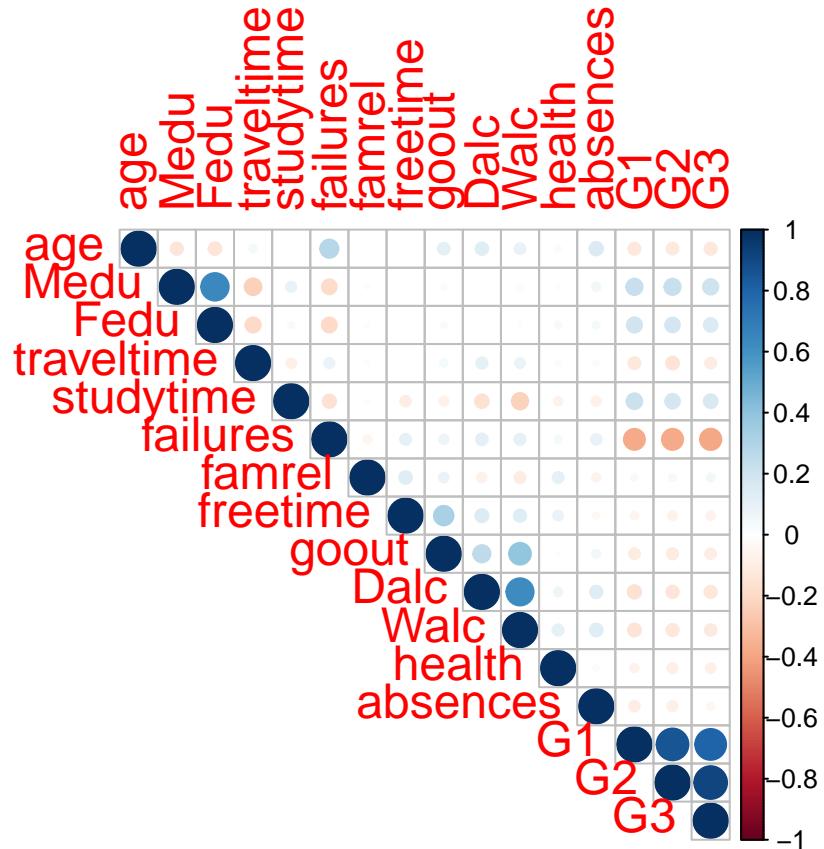
Covariance

```

library(corrplot)

## corrplot 0.95 loaded
numAll = all_data[, sapply(all_data, is.numeric)]
corrMat = cor(numAll)
corrplot(corrMat, type = "upper", number.cex = 1.5, tl.cex = 1.5)

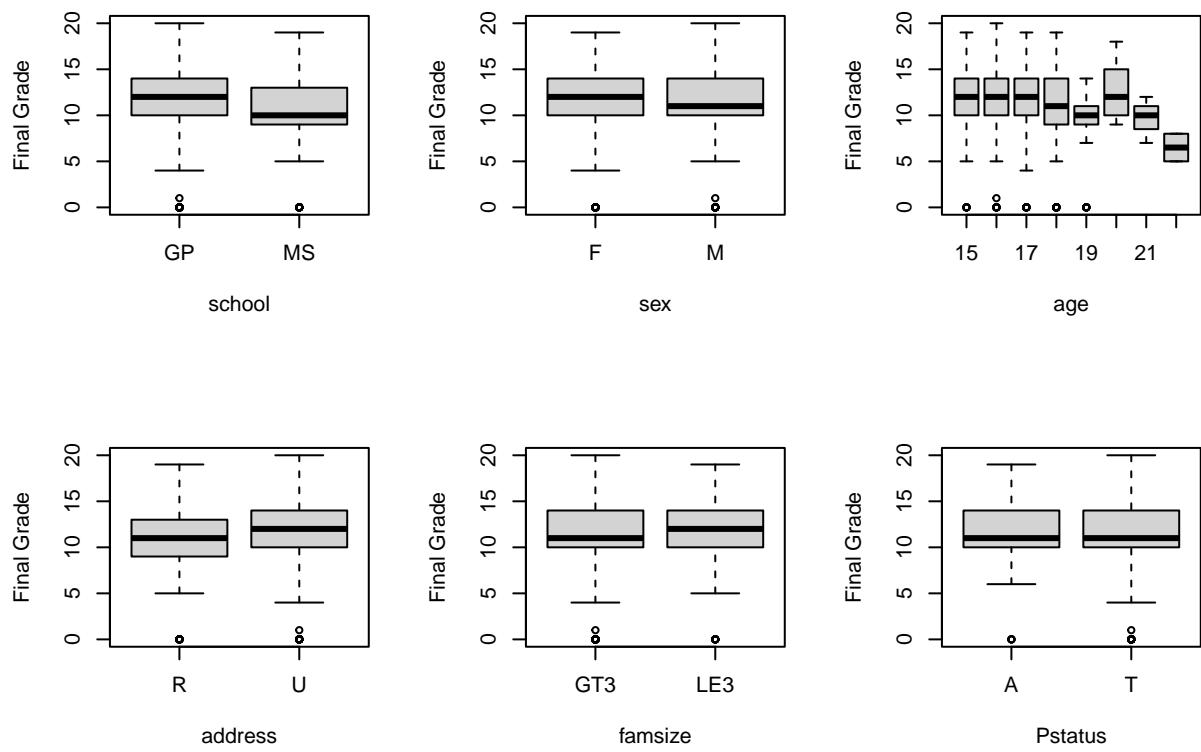
```

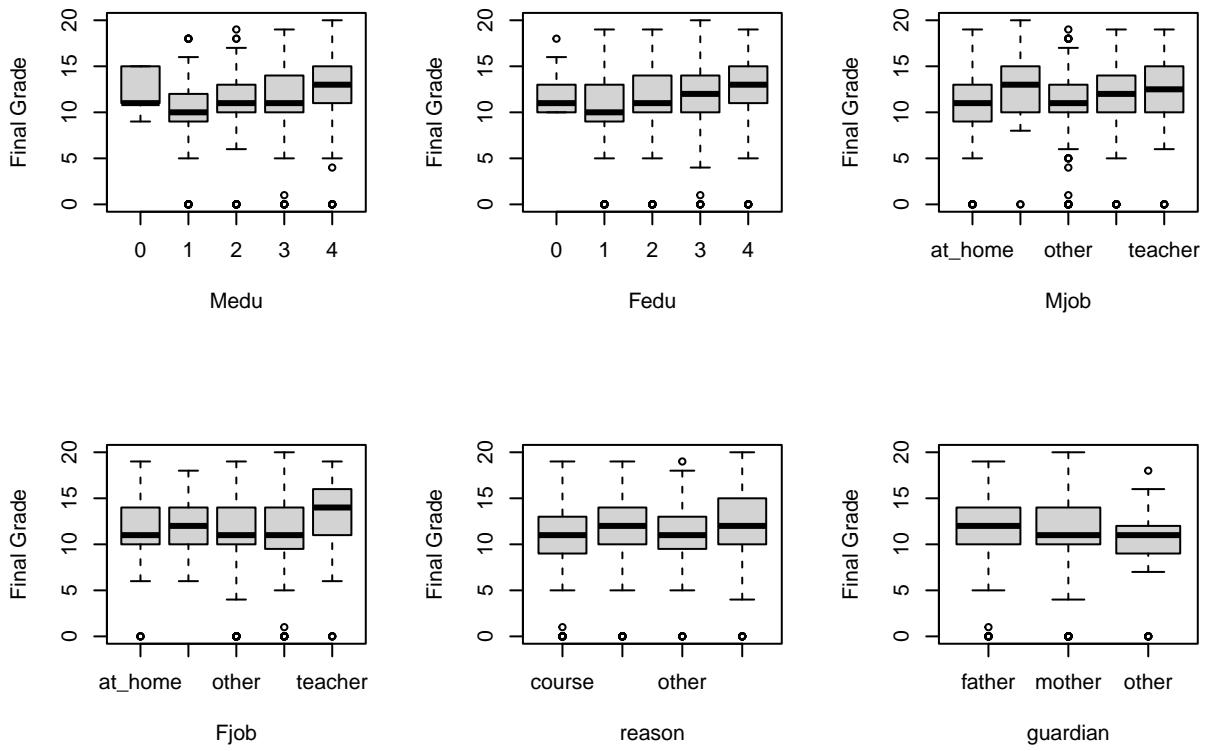


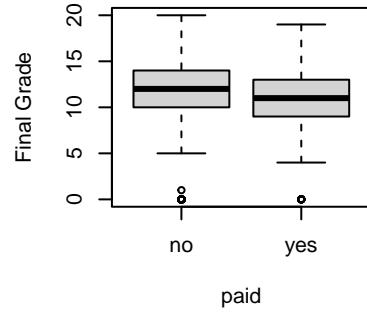
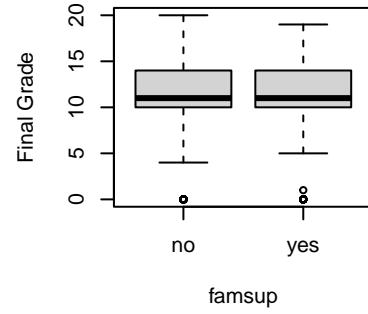
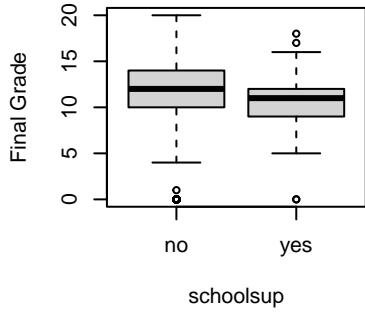
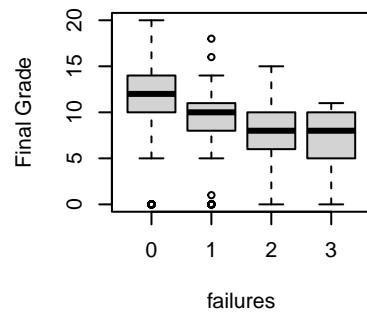
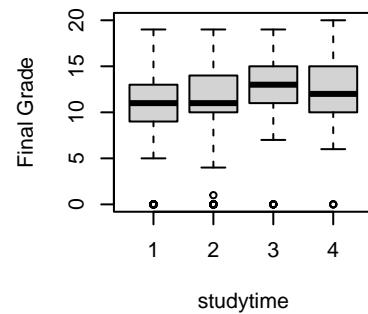
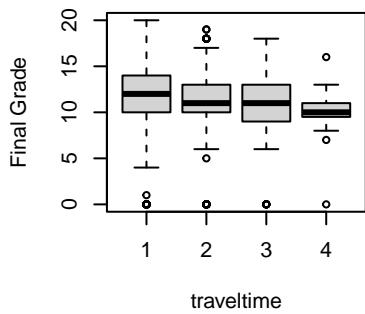
```

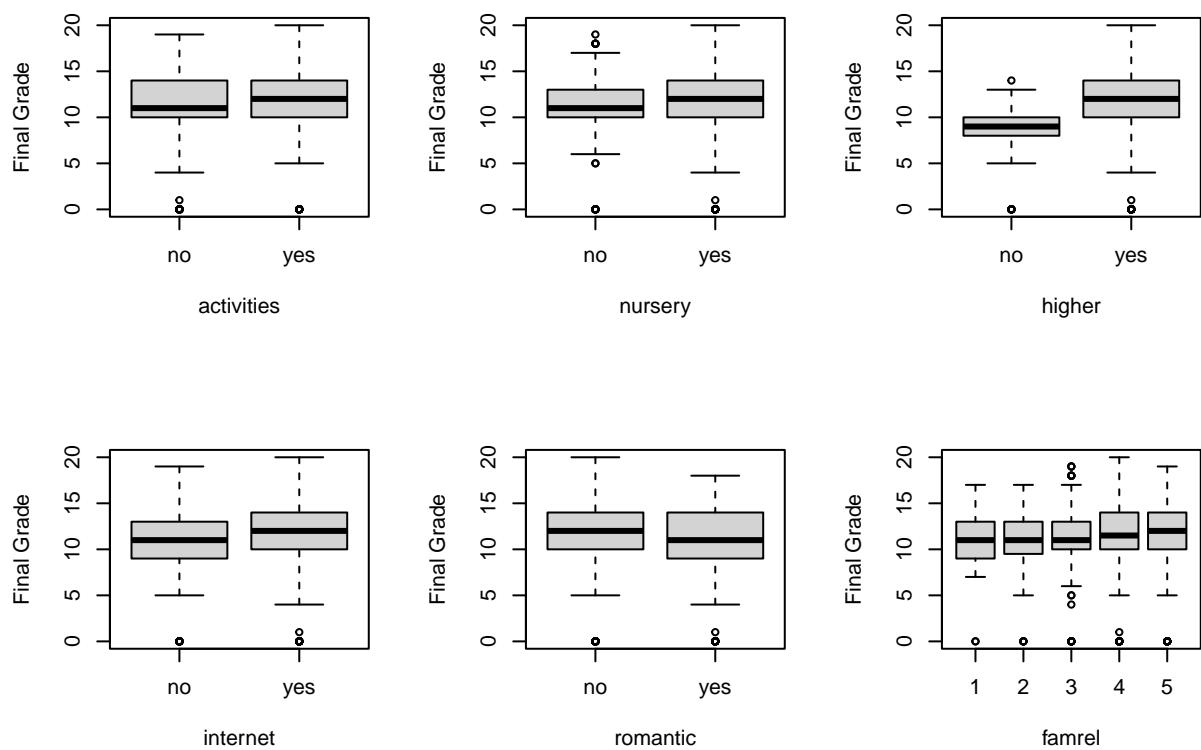
par(mfrow = c(2, 3))
namesInData = names(all_data)
for (i in namesInData){
  if (i != "G3") {
    boxplot(all_data$G3 ~ all_data[, i], xlab = i, ylab = "Final Grade")
  }
}

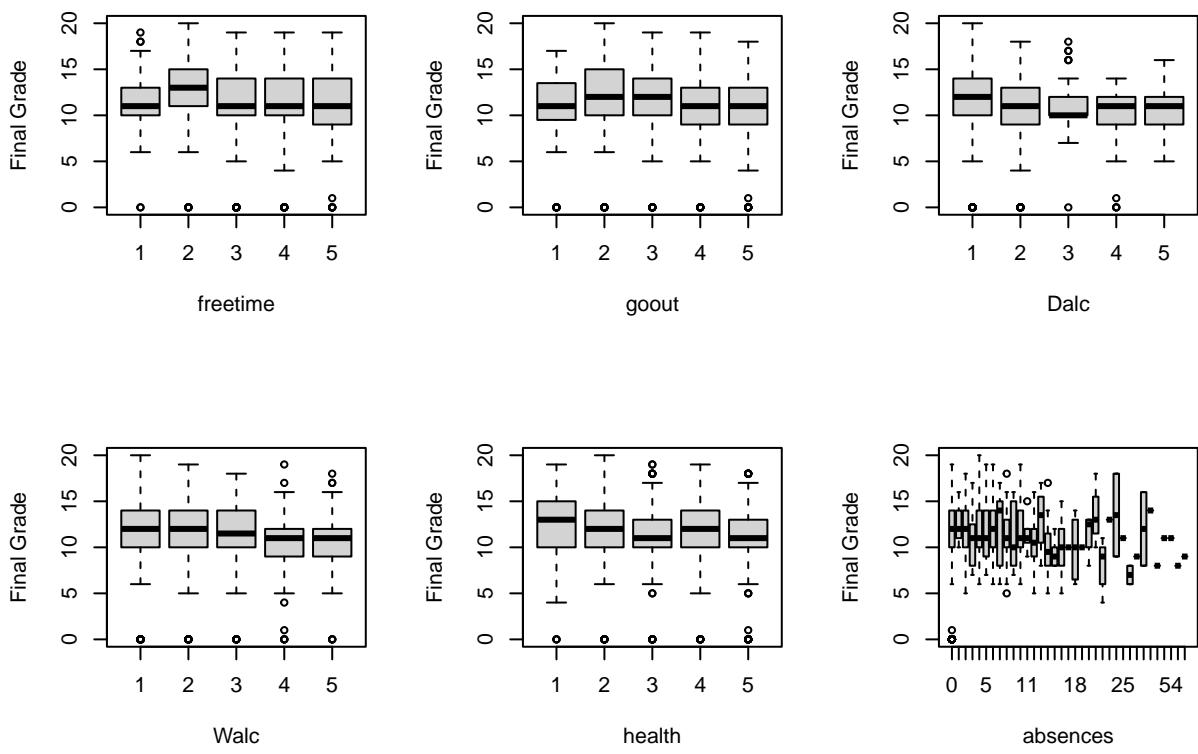
```

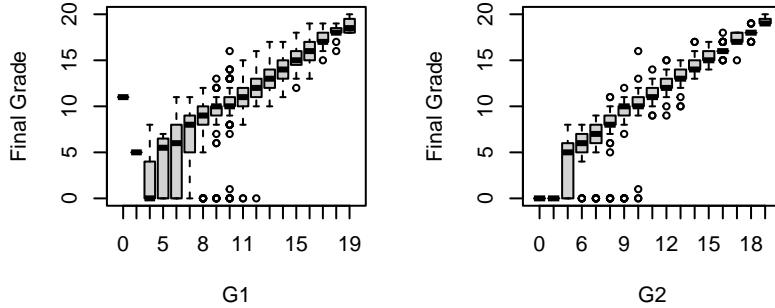






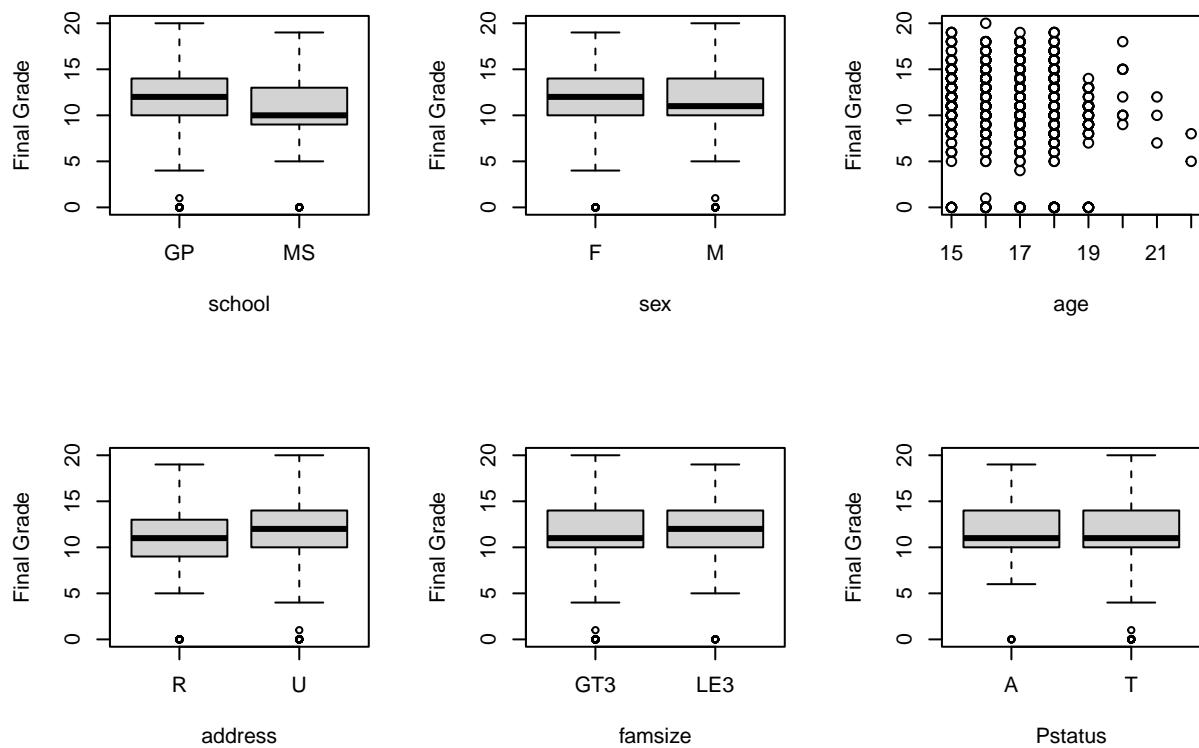


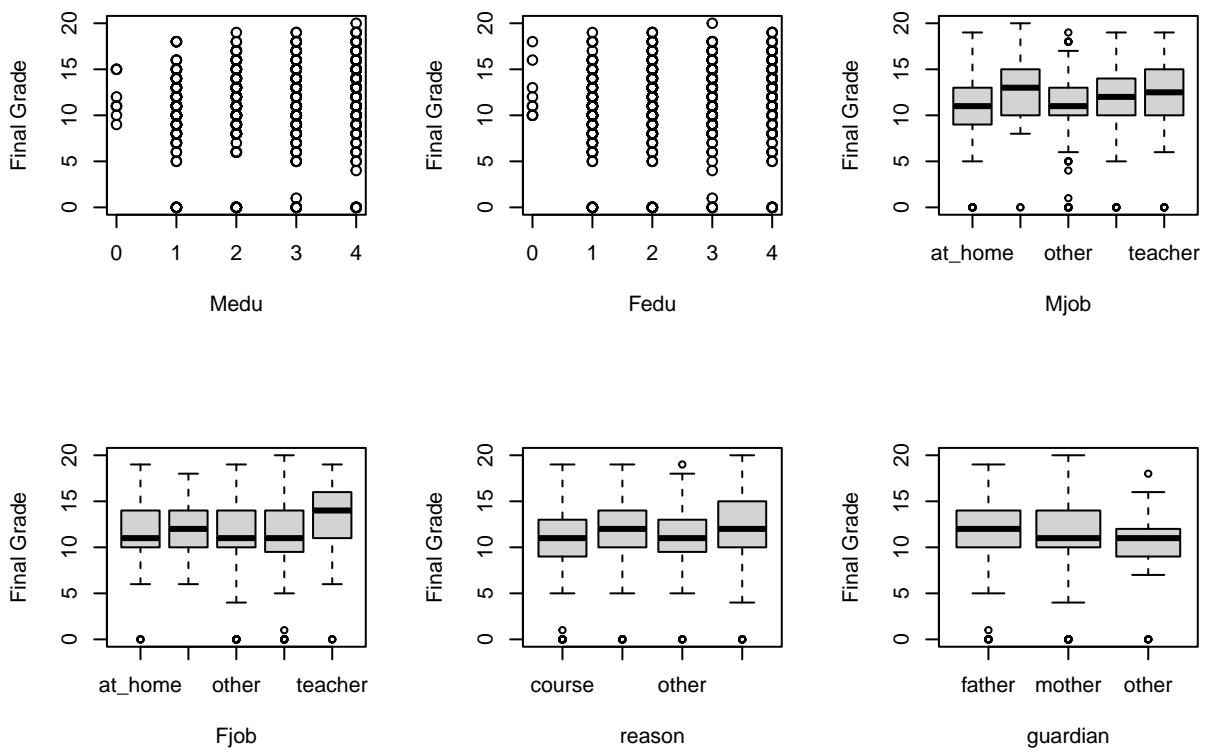


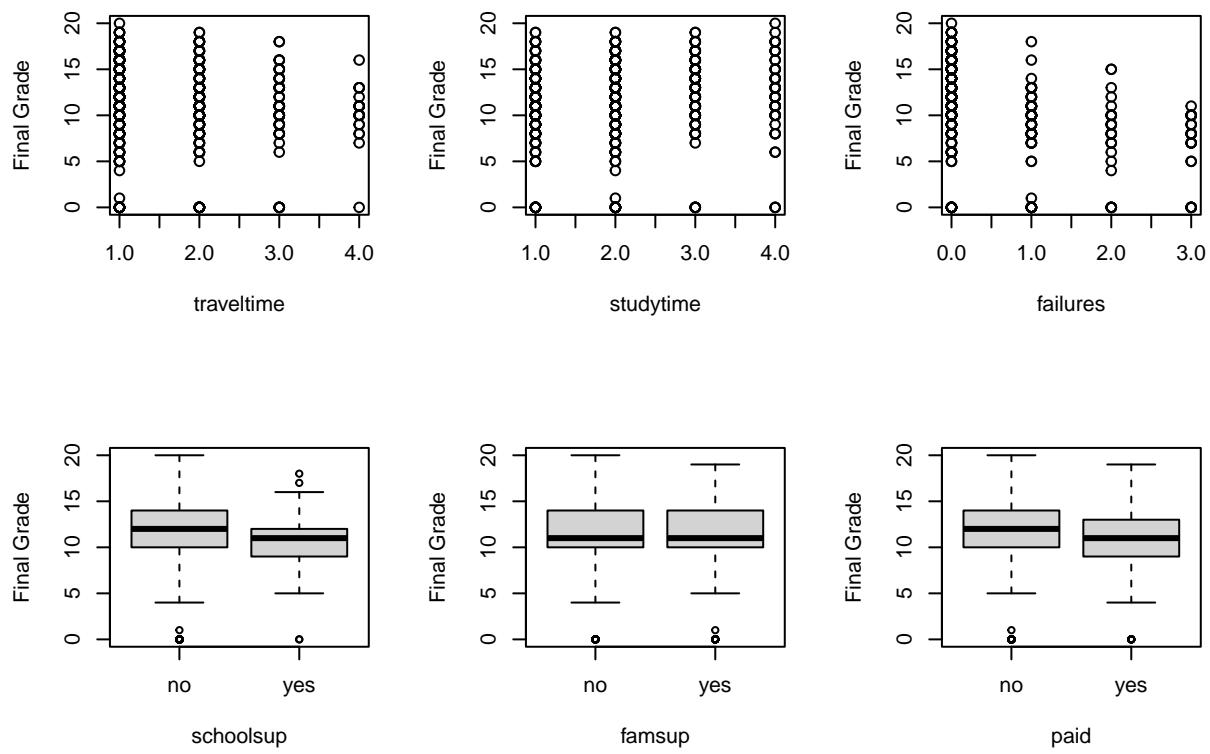


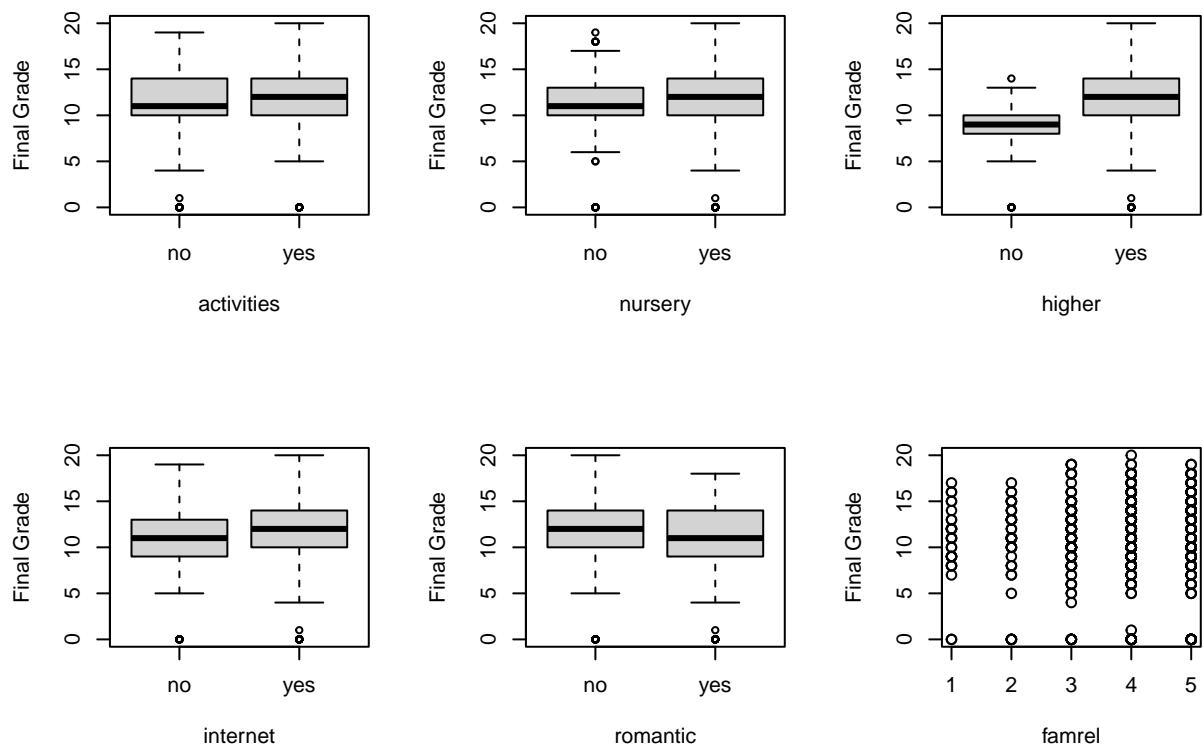
- We see that study time gives a higher final grade along with G1 and G2, which are first and second period grades for the class. We notice the the number of failures the student has had gives and impact for the final grade of there current class.

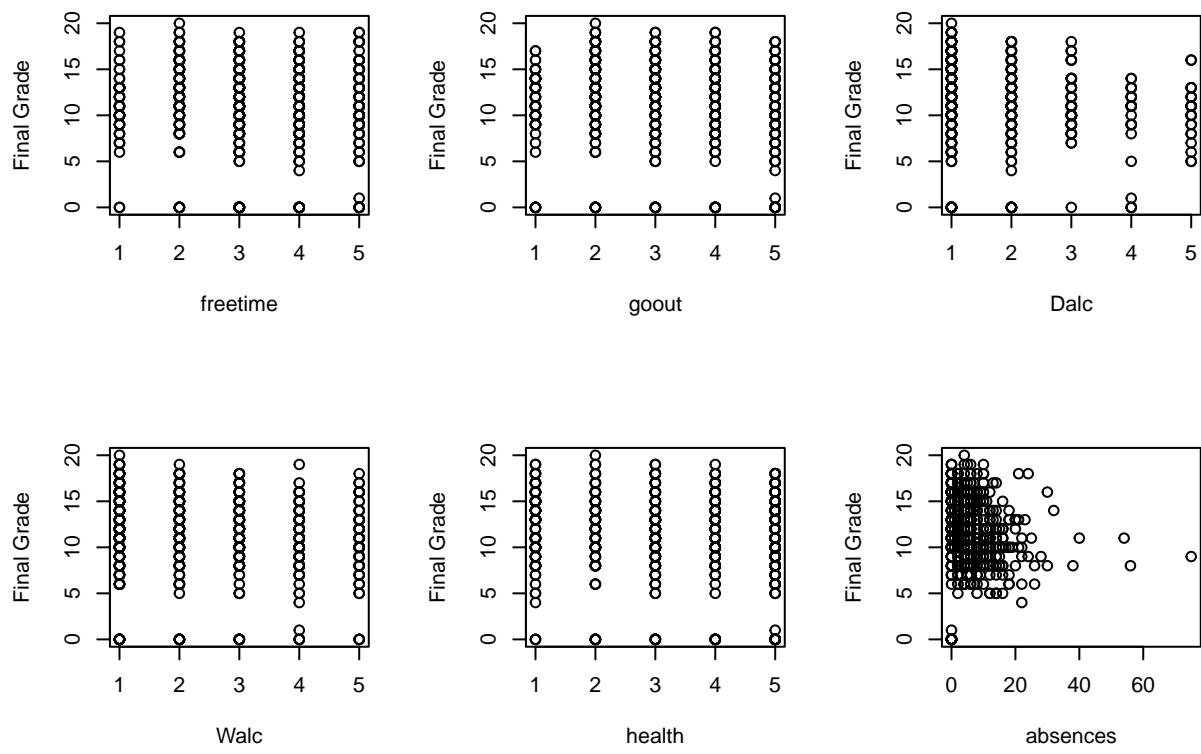
```
par(mfrow = c(2, 3))
namesInData = names(all_data)
for (i in namesInData){
  if (i != "G3") {
    plot(all_data$G3 ~ all_data[, i], xlab = i, ylab = "Final Grade")
  }
}
```

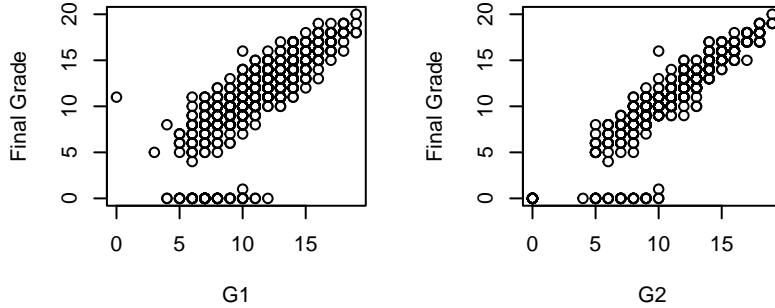












- We need to convert some of the data to factor for example study and travel time since it just classifying a range. In studytime 1 is for less than 2 hours, 2 is for 2 to 5 hours, 3 is for 5 to 10 hours, and 4 is for 10 hours or more. We would have to change these variables so it doesn't mess with our models.

```
nameOfNumericCols = names(all_data)[sapply(all_data, is.numeric)]
nameOfNumericCols
```

```
## [1] "age"         "Medu"        "Fedu"        "traveltime"   "studytime"
## [6] "failures"    "famrel"      "freetime"     "goout"       "Dalc"
## [11] "Walc"        "health"      "absences"    "G1"          "G2"
## [16] "G3"
```

- These are the values that are considered numeric, but we need to change the ones that are just numeric values to define a category. We could change the values for health, Walc, Dalc, goout, and freetime to factors since it just stating a category.

```
set.seed(1)
n = nrow(all_data)
nFeatures= ncol(all_data) - 1
percentTrain = .7
nTrain = n*percentTrain
trainIdx = sample(1:n, nTrain)
trainData = all_data[trainIdx, ]
allLinearFit = lm(G3 ~ ., trainData)
testData = all_data[-trainIdx,]
summary(allLinearFit)
```

```
##
## Call:
```

```

## lm(formula = G3 ~ ., data = trainData)
##
## Residuals:
##   Min     1Q Median     3Q    Max 
## -9.5947 -0.4877  0.1208  0.8346  5.2967 
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.225640  1.170716 -0.193  0.84722  
## schoolMS      0.116010  0.167918  0.691  0.48988  
## sexM          0.059368  0.145215  0.409  0.68279  
## age           -0.007457  0.058443 -0.128  0.89851  
## addressU       0.172259  0.153825  1.120  0.26317  
## famsizeLE3     -0.078917  0.144376 -0.547  0.58483  
## PstatusT      -0.333724  0.213764 -1.561  0.11894  
## Medu          0.060033  0.092483  0.649  0.51647  
## Fedu          -0.097227  0.082065 -1.185  0.23653  
## Mjobhealth    -0.009827  0.307394 -0.032  0.97451  
## Mjobother     -0.075830  0.185939 -0.408  0.68353  
## Mjobservices   0.103891  0.220099  0.472  0.63706  
## Mjobteacher    0.155882  0.295804  0.527  0.59838  
## Fjobhealth    -0.353970  0.448716 -0.789  0.43047  
## Fjobother     -0.363086  0.269014 -1.350  0.17756  
## Fjobservices   -0.636535  0.281288 -2.263  0.02395 *  
## Fjobteacher    -0.657323  0.370821 -1.773  0.07674 .  
## reasonhome    -0.155302  0.165694 -0.937  0.34894  
## reasonother    -0.273730  0.220905 -1.239  0.21572  
## reasonreputation -0.179130  0.169080 -1.059  0.28977  
## guardianmother -0.103299  0.152419 -0.678  0.49817  
## guardianother   -0.090838  0.285108 -0.319  0.75012  
## traveltimes    0.184141  0.091817  2.006  0.04530 *  
## studytime      -0.001513  0.081798 -0.019  0.98524  
## failures        -0.352703  0.110436 -3.194  0.00147 **  
## schoolsupyes   0.161154  0.207622  0.776  0.43790  
## famsupyes      0.304814  0.134382  2.268  0.02362 *  
## paidyes         -0.360163  0.162406 -2.218  0.02690 *  
## activitiesyes  -0.157399  0.129481 -1.216  0.22455  
## nurseryyes     -0.073752  0.160084 -0.461  0.64515  
## higheryes       0.046221  0.239576  0.193  0.84707  
## internetyes    -0.005898  0.166950 -0.035  0.97183  
## romanticyes    -0.081575  0.133812 -0.610  0.54231  
## famrel          0.065754  0.069494  0.946  0.34438  
## freetime        0.030135  0.065734  0.458  0.64679  
## goout           -0.056127  0.063422 -0.885  0.37647  
## Dalc            -0.061533  0.091475 -0.673  0.50138  
## Walc            0.066053  0.072318  0.913  0.36137  
## health          -0.049724  0.044910 -1.107  0.26860  
## absences        0.026853  0.010356  2.593  0.00971 **  
## G1              0.109881  0.040467  2.715  0.00679 **  
## G2              0.970555  0.035448 27.380 < 2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
## 
## Residual standard error: 1.637 on 688 degrees of freedom

```

```

## Multiple R-squared:  0.8373, Adjusted R-squared:  0.8276
## F-statistic: 86.36 on 41 and 688 DF,  p-value: < 2.2e-16
preds = predict(allLinearFit, testData)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.454816

linearFit = lm(G3 ~ G1 + G2 + studytime + failures, trainData)
summary(linearFit)

##
## Call:
## lm(formula = G3 ~ G1 + G2 + studytime + failures, data = trainData)
##
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -9.7815 -0.3544  0.0070  0.8234  5.8788 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.55278   0.28167 -1.962  0.05009 .  
## G1          0.09832   0.03900  2.521  0.01191 *  
## G2          0.97226   0.03475 27.980 < 2e-16 *** 
## studytime   -0.03187   0.07539 -0.423  0.67261    
## failures    -0.30785   0.09722 -3.166  0.00161 ** 
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.644 on 725 degrees of freedom
## Multiple R-squared:  0.8269, Adjusted R-squared:  0.826 
## F-statistic: 866 on 4 and 725 DF,  p-value: < 2.2e-16
preds = predict(linearFit, testData)
sqrt(mean((testData$G3 - preds)^2))

```

[1] 1.419153

-best Subset selection

```

#install.packages("leaps")
library(leaps)
regfit.full = regsubsets(G3 ~ ., trainData, nvmax = nFeatures)
regfit.summary = summary(regfit.full)
regfit.summary

##
## Subset selection object
## Call: regsubsets.formula(G3 ~ ., trainData, nvmax = nFeatures)
## 41 Variables (and intercept)
##          Forced in Forced out
## schoolMS           FALSE      FALSE
## sexM               FALSE      FALSE
## age                FALSE      FALSE
## addressU           FALSE      FALSE
## famsizeLE3          FALSE      FALSE
## PstatusT            FALSE      FALSE
## Medu               FALSE      FALSE

```

```

## Fedu           FALSE    FALSE
## Mjobhealth    FALSE    FALSE
## Mjobother     FALSE    FALSE
## Mjobservices   FALSE    FALSE
## Mjobteacher    FALSE    FALSE
## Fjobhealth    FALSE    FALSE
## Fjobother     FALSE    FALSE
## Fjobservices   FALSE    FALSE
## Fjobteacher    FALSE    FALSE
## reasonhome    FALSE    FALSE
## reasonother    FALSE    FALSE
## reasonreputation FALSE    FALSE
## guardianmother FALSE    FALSE
## guardianother   FALSE    FALSE
## travelttime   FALSE    FALSE
## studytime     FALSE    FALSE
## failures       FALSE    FALSE
## schoolsupyes   FALSE    FALSE
## famsupyes      FALSE    FALSE
## paidyes        FALSE    FALSE
## activitiesyes  FALSE    FALSE
## nurseryyes     FALSE    FALSE
## higheryes      FALSE    FALSE
## internetyes    FALSE    FALSE
## romanticyes    FALSE    FALSE
## famrel         FALSE    FALSE
## freetime        FALSE    FALSE
## goout          FALSE    FALSE
## Dalc            FALSE    FALSE
## Walc            FALSE    FALSE
## health          FALSE    FALSE
## absences        FALSE    FALSE
## G1              FALSE    FALSE
## G2              FALSE    FALSE

## 1 subsets of each size up to 32
## Selection Algorithm: exhaustive
##          schoolMS sexM age addressU famsizeLE3 PstatusT Medu Fedu Mjobhealth
## 1  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 2  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 3  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 4  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 5  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 6  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 7  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 8  ( 1 )   " "   " "   " "   " "   " "   " "   " "   " "
## 9  ( 1 )   " "   " "   " "   " "   " "   "*"  " "   " "
## 10 ( 1 )   " "   " "   " "   " "   " "   "*"  " "   " "
## 11 ( 1 )   " "   " "   " "   " "   " "   "*"  " "   " "
## 12 ( 1 )   " "   " "   " "   "*"  " "   "*"  " "   " "
## 13 ( 1 )   " "   " "   " "   "*"  " "   "*"  " "   " "
## 14 ( 1 )   " "   " "   " "   "*"  " "   "*"  " "   " "
## 15 ( 1 )   " "   " "   " "   "*"  " "   "*"  " "   " "
## 16 ( 1 )   " "   " "   " "   "*"  " "   "*"  " "   " "
## 17 ( 1 )   " "   " "   " "   "*"  " "   "*"  " "   " "

```

```

## 18 ( 1 ) " " " " " * " " " " " * "
## 19 ( 1 ) " " " " " * " " " " " * "
## 20 ( 1 ) " " " " " * " " " " " * "
## 21 ( 1 ) " " " " " * " " " " " * "
## 22 ( 1 ) " " " " " * " " " " " * "
## 23 ( 1 ) " " " " " * " " " " " * "
## 24 ( 1 ) " " " " " * " " " " " * "
## 25 ( 1 ) " " " " " * " " " " " * "
## 26 ( 1 ) " " " " " * " " " " " * "
## 27 ( 1 ) " * " " " " * " " " " " * "
## 28 ( 1 ) " * " " " " * " " " " " * "
## 29 ( 1 ) " * " " " " * " " " " " * "
## 30 ( 1 ) " * " " " " * " " " " " * "
## 31 ( 1 ) " * " " " " * " " * " " * "
## 32 ( 1 ) " * " " " " " * " " * " " * "
##
## Mjobother Mjobservices Mjobteacher Fjobhealth Fjobother Fjobservices
## 1 ( 1 ) " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " "
## 4 ( 1 ) " " " " " " " " " "
## 5 ( 1 ) " " " " " " " " " "
## 6 ( 1 ) " " " " " " " " " "
## 7 ( 1 ) " " " " " " " " " "
## 8 ( 1 ) " " " " " " " " " * "
## 9 ( 1 ) " " " " " " " " " * "
## 10 ( 1 ) " " " " " " " " " * "
## 11 ( 1 ) " " " " " " " " " * "
## 12 ( 1 ) " " " " " " " " " * "
## 13 ( 1 ) " " " " " " " " * " " * "
## 14 ( 1 ) " " " " " " * " " * "
## 15 ( 1 ) " " " " " " * " " * "
## 16 ( 1 ) " " " " " " * " " * "
## 17 ( 1 ) " * " " " " " * " " * "
## 18 ( 1 ) " * " " " " " * " " * "
## 19 ( 1 ) " " " " " " * " " * "
## 20 ( 1 ) " * " " " " " * " " * "
## 21 ( 1 ) " * " " " " " * " " * "
## 22 ( 1 ) " " " " " " * " " * "
## 23 ( 1 ) " * " " " " " * " " * "
## 24 ( 1 ) " " " " " " * " " * "
## 25 ( 1 ) " * " " " " " * " " * "
## 26 ( 1 ) " * " " " " " * " " * "
## 27 ( 1 ) " * " " " " " * " " * "
## 28 ( 1 ) " * " " " " " * " " * "
## 29 ( 1 ) " " " " * " " * " " * "
## 30 ( 1 ) " " " " * " " * " " * "
## 31 ( 1 ) " " " " * " " * " " * "
## 32 ( 1 ) " " " " * " " * " " * "
##
## Fjobteacher reasonhome reasonother reasonreputation guardianmother
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " " " " " " "
## 3 ( 1 ) " " " " " " " "
## 4 ( 1 ) " " " " " " " "
## 5 ( 1 ) " " " " " " " "

```

```

## 6  ( 1 )   " "      " "      " "      " "
## 7  ( 1 )   " "      " "      " "      " "
## 8  ( 1 )   " "      " "      " "      " "
## 9  ( 1 )   " "      " "      " "      " "
## 10 ( 1 )   "*"     " "      " "      " "
## 11 ( 1 )   "*"     " "      " "      " "
## 12 ( 1 )   "*"     " "      " "      " "
## 13 ( 1 )   "*"     " "      " "      " "
## 14 ( 1 )   "*"     " "      " "      " "
## 15 ( 1 )   "*"     " "      " "      " "
## 16 ( 1 )   "*"     " "      "*"     " "
## 17 ( 1 )   "*"     " "      "*"     " "
## 18 ( 1 )   "*"     " "      "*"     " "
## 19 ( 1 )   "*"     "*"     "*"     " "
## 20 ( 1 )   "*"     "*"     "*"     " "
## 21 ( 1 )   "*"     "*"     "*"     " "
## 22 ( 1 )   "*"     "*"     "*"     " "
## 23 ( 1 )   "*"     "*"     "*"     " "
## 24 ( 1 )   "*"     "*"     "*"     " "
## 25 ( 1 )   "*"     "*"     "*"     " "
## 26 ( 1 )   "*"     "*"     "*"     " "
## 27 ( 1 )   "*"     "*"     "*"     " "
## 28 ( 1 )   "*"     "*"     "*"     " "
## 29 ( 1 )   "*"     "*"     "*"     " "
## 30 ( 1 )   "*"     "*"     "*"     "*"
## 31 ( 1 )   "*"     "*"     "*"     "*"
## 32 ( 1 )   "*"     "*"     "*"     "*"

##          guardianother travelttime studytime failures schoolsupyes famsupyes
## 1  ( 1 )   " "      " "      " "      " "      " "      " "
## 2  ( 1 )   " "      " "      " "      "*"     " "      " "
## 3  ( 1 )   " "      " "      " "      "*"     " "      " "
## 4  ( 1 )   " "      " "      " "      "*"     " "      " "
## 5  ( 1 )   " "      " "      " "      "*"     " "      " "
## 6  ( 1 )   " "      " "      " "      "*"     " "      "*"
## 7  ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 8  ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 9  ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 10 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 11 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 12 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 13 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 14 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 15 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 16 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 17 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 18 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 19 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 20 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 21 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 22 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 23 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 24 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 25 ( 1 )   " "      "*"     " "      "*"     " "      "*"
## 26 ( 1 )   " "      "*"     " "      "*"     " "      "*"

```

```

## 27 ( 1 ) " "      "*"      " "      "*"      "*"      "*"
## 28 ( 1 ) " "      "*"      " "      "*"      "*"      "*"
## 29 ( 1 ) " "      "*"      " "      "*"      "*"      "*"
## 30 ( 1 ) " "      "*"      " "      "*"      "*"      "*"
## 31 ( 1 ) " "      "*"      " "      "*"      "*"      "*"
## 32 ( 1 ) " "      "*"      " "      "*"      "*"      "*"
##          paidyes activtiesyes nurseryyes higheryes internetyes romanticyes
## 1 ( 1 ) " "      " "      " "      " "      " "      " "
## 2 ( 1 ) " "      " "      " "      " "      " "      " "
## 3 ( 1 ) " "      " "      " "      " "      " "      " "
## 4 ( 1 ) " "      " "      " "      " "      " "      " "
## 5 ( 1 ) "*"      " "      " "      " "      " "      " "
## 6 ( 1 ) "*"      " "      " "      " "      " "      " "
## 7 ( 1 ) "*"      " "      " "      " "      " "      " "
## 8 ( 1 ) "*"      " "      " "      " "      " "      " "
## 9 ( 1 ) "*"      " "      " "      " "      " "      " "
## 10 ( 1 ) "*"     " "      " "      " "      " "      " "
## 11 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 12 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 13 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 14 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 15 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 16 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 17 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 18 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 19 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 20 ( 1 ) "*"     "*"     " "      " "      " "      " "
## 21 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 22 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 23 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 24 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 25 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 26 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 27 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 28 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 29 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 30 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 31 ( 1 ) "*"     "*"     " "      " "      " "      "*"
## 32 ( 1 ) "*"     "*"     " "      " "      " "      "*"
##          famrel freetime goout Dalc Walc health absences G1  G2
## 1 ( 1 ) " "      " "      " "      " "      " "      " "      " "      "*"
## 2 ( 1 ) " "      " "      " "      " "      " "      " "      " "      "*"
## 3 ( 1 ) " "      " "      " "      " "      " "      "*"      "*"      "*"
## 4 ( 1 ) " "      " "      " "      " "      " "      "*"      "*"      "*"
## 5 ( 1 ) " "      " "      " "      " "      " "      "*"      "*"      "*"
## 6 ( 1 ) " "      " "      " "      " "      " "      "*"      "*"      "*"
## 7 ( 1 ) " "      " "      " "      " "      " "      "*"      "*"      "*"
## 8 ( 1 ) " "      " "      " "      " "      " "      "*"      "*"      "*"
## 9 ( 1 ) " "      " "      " "      " "      " "      "*"      "*"      "*"
## 10 ( 1 ) " "     " "      " "      " "      " "      "*"      "*"      "*"
## 11 ( 1 ) " "     " "      " "      " "      " "      "*"      "*"      "*"
## 12 ( 1 ) " "     " "      " "      " "      " "      "*"      "*"      "*"
## 13 ( 1 ) " "     " "      " "      " "      " "      "*"      "*"      "*"
## 14 ( 1 ) " "     " "      " "      " "      " "      "*"      "*"      "*"

```

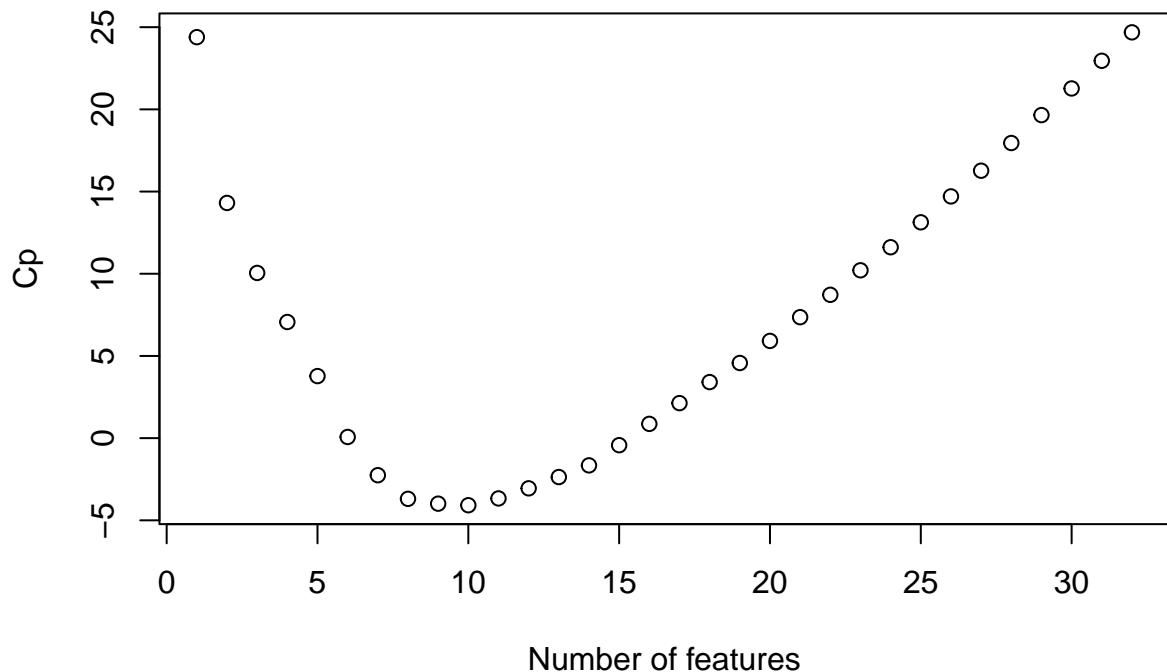
```

## 15  ( 1 ) " "   " "   " "   " "   " "   " "   " "   "*"   "*"   "*"
## 16  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 17  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 18  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 19  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 20  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 21  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 22  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 23  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 24  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 25  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 26  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 27  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 28  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 29  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 30  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 31  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"
## 32  ( 1 ) "*"  " "   " "   " "   " "   " "   "*"   "*"   "*"   "*"

regfit.adjr2 = regfit.summary$adjr2
plot(regfit.summary$cp, xlab = "Number of features", ylab = "Cp", main = "Cp vs Number of features")

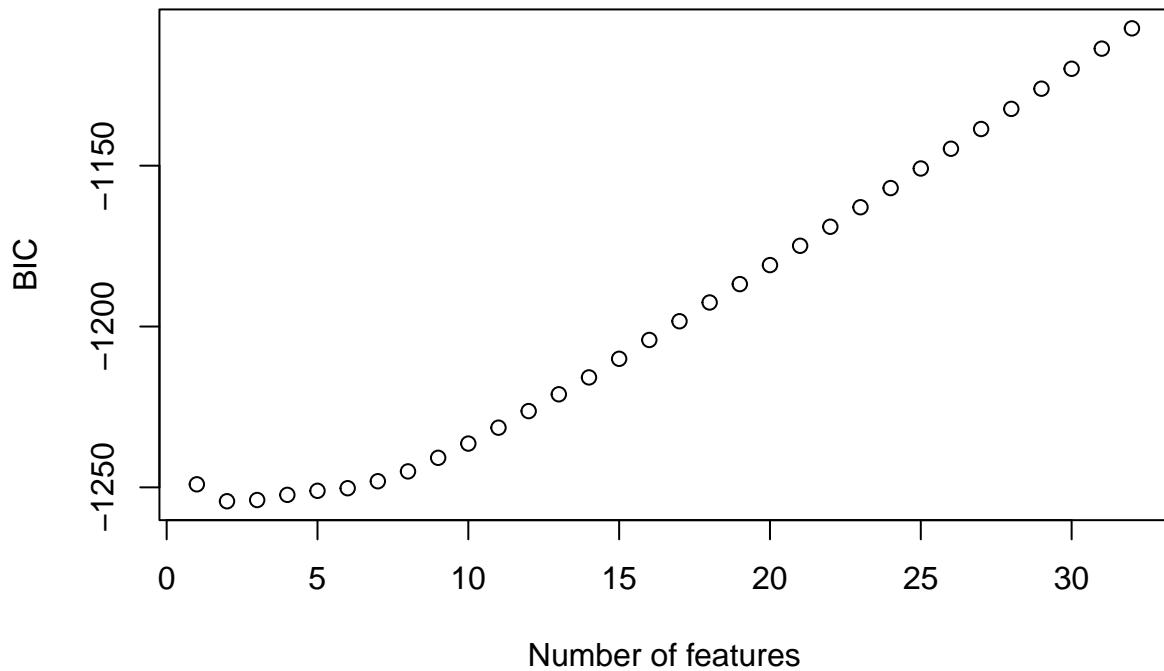
```

Cp vs Number of features



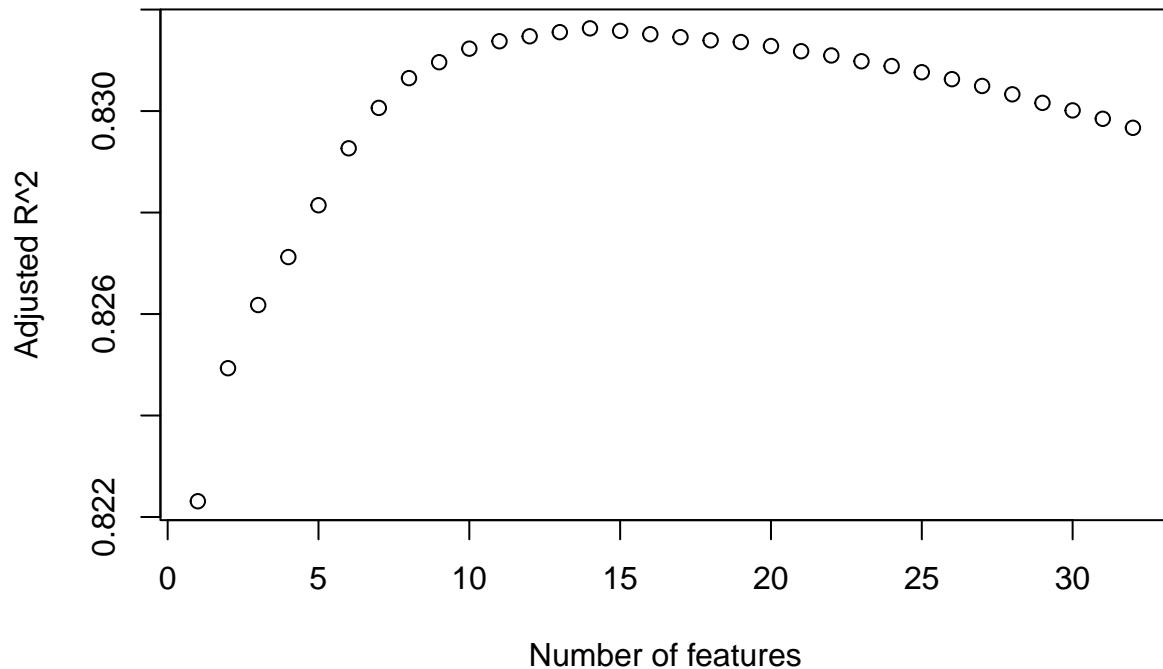
```
plot(regfit.summary$bic, xlab = "Number of features", ylab = "BIC", main = "BIC vs Number of features")
```

BIC vs Number of features



```
plot(regfit.adjr2, xlab = "Number of features", ylab = "Adjusted R^2", main = "Adjusted R^2 vs Number of features")
```

Adjusted R² vs Number of features



```

which.max(regfit.adjr2)
## [1] 14
which.min(regfit.summary$cp)
## [1] 10
which.min(regfit.summary$bic)
## [1] 2
regfit.adjr2[14]
## [1] 0.83163
coef(regfit.full, id=14)

## (Intercept) addressU PstatusT Fjobhealth Fjobother
## -0.55602684 0.16446304 -0.27599634 -0.47996236 -0.41070208
## Fjobservices Fjobteacher traveltim failures famsupyes
## -0.62719616 -0.74383734 0.20491426 -0.34479517 0.31277034
## paidyes activitiesyes absences G1 G2
## -0.37572559 -0.16167604 0.02279686 0.10364214 0.97581309

coef(regfit.full, id=10)
## (Intercept) PstatusT Fjobservices Fjobteacher traveltim failures
## -0.81012477 -0.31331663 -0.25521576 -0.35933059 0.17347400 -0.33761632
## famsupyes paidyes absences G1 G2

```

```

##   0.32706672 -0.38763466  0.02279829  0.10161936  0.97692118
coef(regfit.full, id=2)

## (Intercept) failures          G2
## -0.3073145 -0.3336223  1.0433859

• The best model according to the training data is with 14 variables. The variables are address, Pstatus, Fjob, travelttime, failures, famsup, paid, activities, absences, G1, and G2. This is not a total of 14 because it has dummy variables within Fjob, and famsup. We can try to make a model with these variables. This turns out to be worse then forward selection so we should try to find the model that has the best CV.

bestSub14fit = lm(G3 ~ address + Pstatus + Fjob + travelttime + failures + famsup + paid + activities + a
summary(bestSub14fit)

##
## Call:
## lm(formula = G3 ~ address + Pstatus + Fjob + travelttime + failures +
##     famsup + paid + activities + absences + G1 + G2, data = trainData)
##
## Residuals:
##    Min      1Q  Median      3Q      Max
## -9.5110 -0.4667  0.1350  0.8348  5.5295
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.55603   0.44635 -1.246 0.213271
## addressU     0.16446   0.14031  1.172 0.241536
## PstatusT    -0.27600   0.19825 -1.392 0.164306
## Fjobhealth  -0.47996   0.41776 -1.149 0.250980
## Fjobother   -0.41070   0.25690 -1.599 0.110337
## Fjobservices -0.62720   0.26609 -2.357 0.018686 *
## Fjobteacher  -0.74384   0.33711 -2.207 0.027665 *
## travelttime   0.20491   0.08581  2.388 0.017201 *
## failures     -0.34480   0.09614 -3.587 0.000358 ***
## famsupyes    0.31277   0.12724  2.458 0.014202 *
## paidyes      -0.37573   0.15476 -2.428 0.015437 *
## activitiesyes -0.16168   0.12099 -1.336 0.181879
## absences      0.02280   0.00950  2.400 0.016660 *
## G1            0.10364   0.03848  2.693 0.007245 **
## G2            0.97581   0.03447 28.313 < 2e-16 ***
##
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.617 on 715 degrees of freedom
## Multiple R-squared:  0.8349, Adjusted R-squared:  0.8316
## F-statistic: 258.2 on 14 and 715 DF,  p-value: < 2.2e-16

preds = predict(bestSub14fit, testData)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.440578

bestSub10fit = lm(G3 ~ Pstatus + Fjob + travelttime + failures + paid + famsup + absences + G1 + G2, tra
summary(bestSub10fit)

```

```

## 
## Call:
## lm(formula = G3 ~ Pstatus + Fjob + travelttime + failures + paid +
##      famsup + absences + G1 + G2, data = trainData)
##
## Residuals:
##    Min     1Q Median     3Q    Max 
## -9.5557 -0.4579  0.1305  0.8408  5.5619 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.473595  0.419284 -1.130 0.259051    
## PstatusT     -0.318071  0.196979 -1.615 0.106806    
## Fjobhealth   -0.440912  0.417212 -1.057 0.290956    
## Fjobother    -0.393856  0.256652 -1.535 0.125325    
## Fjobservices -0.612581  0.266094 -2.302 0.021614 *  
## Fjobteacher   -0.717692  0.337035 -2.129 0.033559 *  
## travelttime   0.180461  0.082118  2.198 0.028298 *  
## failures      -0.341023  0.096196 -3.545 0.000418 *** 
## paidyes       -0.377849  0.154392 -2.447 0.014630 *  
## famsupyes     0.318290  0.127282  2.501 0.012618 *  
## absences       0.022881  0.009506  2.407 0.016331 *  
## G1              0.101446  0.038498  2.635 0.008593 **  
## G2              0.978772  0.034441 28.419 < 2e-16 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
##
## Residual standard error: 1.619 on 717 degrees of freedom 
## Multiple R-squared:  0.8341, Adjusted R-squared:  0.8313 
## F-statistic: 300.4 on 12 and 717 DF,  p-value: < 2.2e-16 

preds = predict(bestSub10fit, testData)
sqrt(mean((testData$G3 - preds)^2))

```

[1] 1.433748

- Using validation set for best subset selection

```

set.seed(1)
percentVal = .2
nValidation = percentVal*n
valD = sample(1:nTrain, nValidation)
subSetVal = trainData[valD,]
subSetTrain = trainData[-valD,]

```

```
regfit.best.validate = regsubsets(G3 ~ ., subSetTrain, nvmax = nFeatures)
```

- the model with the lowest validation error is one that uses 7 variables. They are travelttime, failures, famsup, paid, absences, G1, and G2. Now we should train a model using the full training set and check the testing error.

```

val.mat = model.matrix(G3 ~ ., subSetVal)
val.errors = rep(NA, nFeatures)
for (i in 1:(ncol(trainData) - 1)) {
  coefi = coef(regfit.best.validate, id = i)
  pred = val.mat[, names(coefi)] %*% coefi
  val.errors[i] = mean((subSetVal$G3 - pred)^2)
}

```

```

}

which.min(val.errors)

## [1] 5

coef(regfit.best.validate, which.min(val.errors))

## (Intercept) travelltime failures absences          G1          G2
## -1.07793860  0.23087825 -0.36886219  0.02389852  0.09505954  0.98074559

coef(regfit.full, 5)

## (Intercept) failures paidyes absences          G1          G2
## -0.62551739 -0.33412055 -0.35222242  0.02373891  0.09653667  0.97228413

- This has the lowest RMSE

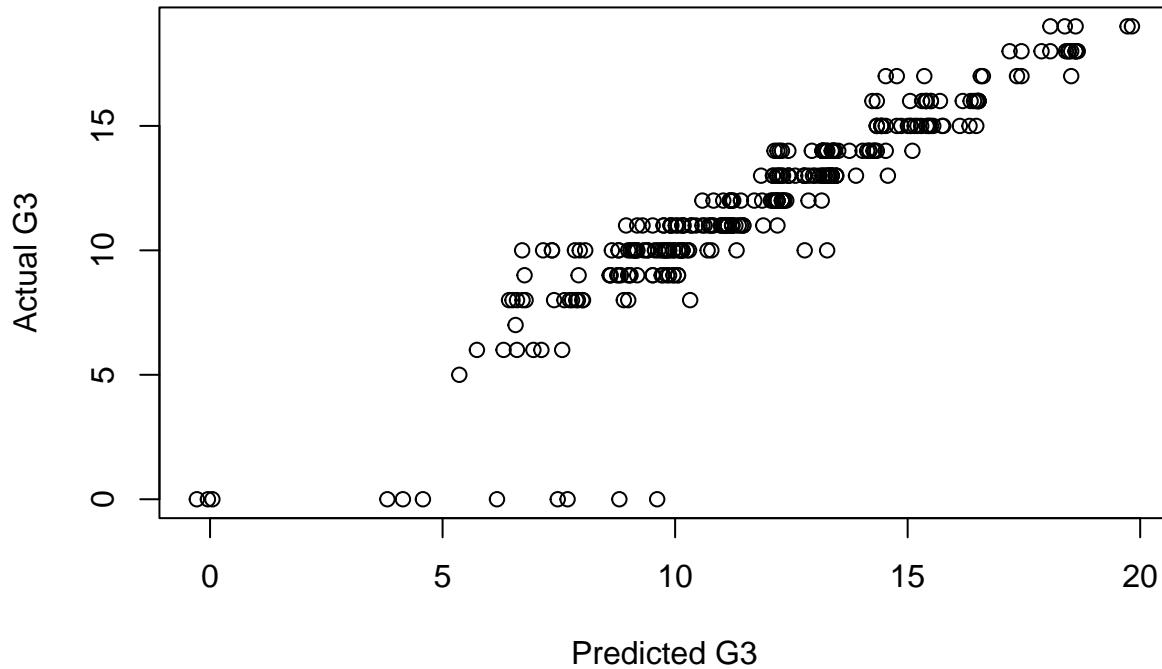
bestSub5Fit = lm(G3 ~ failures + paid + absences + G1 + G2, trainData)
summary(bestSub5Fit)

## Call:
## lm(formula = G3 ~ failures + paid + absences + G1 + G2, data = trainData)
##
## Residuals:
##    Min      1Q  Median      3Q     Max
## -9.7286 -0.3855  0.0801  0.8270  5.9373
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.625517   0.270215 -2.315 0.020898 *
## failures     -0.334121   0.096900 -3.448 0.000597 ***
## paidyes      -0.352222   0.152885 -2.304 0.021514 *
## absences      0.023739   0.009546  2.487 0.013114 *
## G1            0.096537   0.038586  2.502 0.012574 *
## G2            0.972284   0.034538 28.151 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

## Residual standard error: 1.634 on 724 degrees of freedom
## Multiple R-squared:  0.8293, Adjusted R-squared:  0.8281
## F-statistic: 703.6 on 5 and 724 DF,  p-value: < 2.2e-16

pred = predict(bestSub5Fit, testData)
plot(pred, testData$G3, xlab = "Predicted G3", ylab = "Actual G3", main = "Predicted vs Actual Final Gr
```

Predicted vs Actual Final Grades (Best 5 Subset Validation Set)



```
sqrt(mean((testData$G3 - pred)^2))
```

```
## [1] 1.398733
```

- Try doing CV with best subset. Must make prediction function. According to the output we should use 2 features for the lowest mean cv error

```
predict.regsubsets = function(object, newData, id, ...) {
  form = as.formula(object$call[[2]])
  mat = model.matrix(form, newData)
  coefi = coef(object, id = id)
  xvars = names(coefi)
  mat[, xvars] %*% coefi
}
# 10 fold CV
k = 10
set.seed(1)
folds = sample(rep(1:k, length = nTrain))
cv.errors = matrix(NA, k, nFeatures, dimnames = list(NULL, paste(1:nFeatures)))

for (j in 1:k) {
  best.fit = regsubsets(G3 ~ ., trainData[folds != j, ], nvmax = nFeatures)
  for (i in 1:nFeatures) {
    pred = predict(best.fit, trainData[folds == j, ], nvmax = nFeatures, id = i)
    cv.errors[j, i] = mean((trainData$G3[folds == j] - pred)^2)
  }
}
```

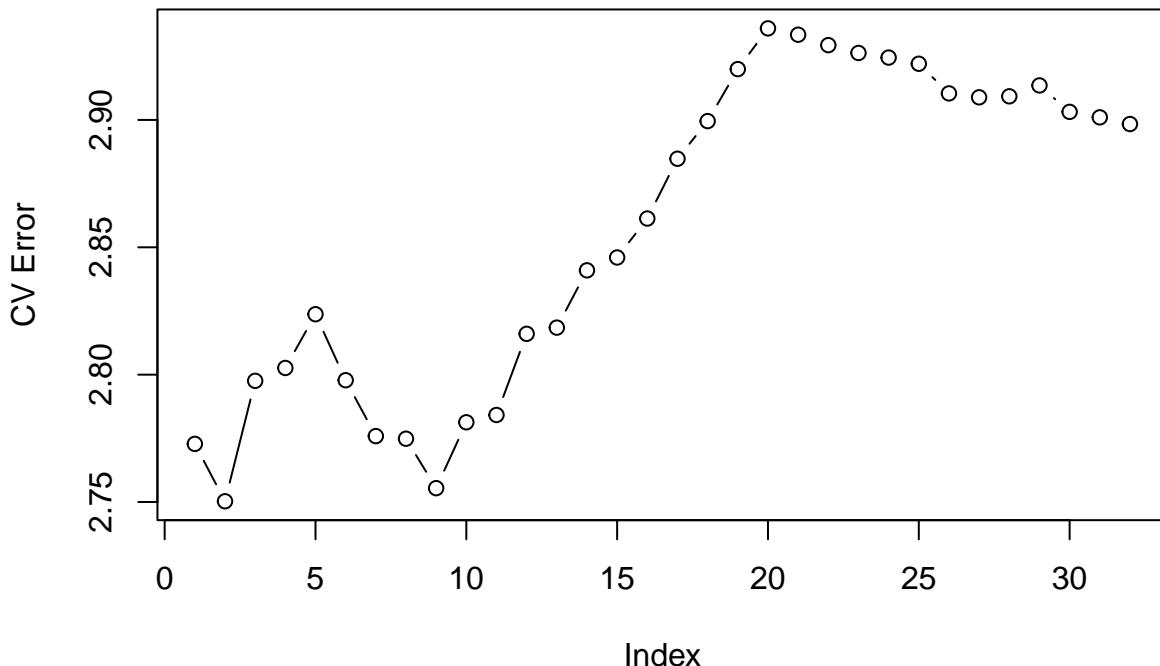
```

mean.cv.errors = apply(cv.errors, 2, mean)
mean.cv.errors

##      1      2      3      4      5      6      7      8
## 2.772813 2.750289 2.797568 2.802625 2.823720 2.797807 2.775868 2.774846
##      9     10     11     12     13     14     15     16
## 2.755433 2.781296 2.784160 2.816028 2.818463 2.840950 2.845994 2.861294
##     17     18     19     20     21     22     23     24
## 2.884765 2.899538 2.919952 2.935950 2.933457 2.929371 2.926291 2.924485
##     25     26     27     28     29     30     31     32
## 2.922064 2.910421 2.908880 2.909273 2.913549 2.903176 2.901019 2.898393
par(mfrow = c(1,1))
plot(mean.cv.errors, type= "b", ylab = "CV Error", main = "CV Error vs number of best subset Features")

```

CV Error vs number of best subset Features



```

which.min(mean.cv.errors)

## 2
## 2

coef(regfit.full, which.min(mean.cv.errors))

## (Intercept) failures          G2
## -0.3073145 -0.3336223   1.0433859

bestSub2fit = lm(G3 ~ failures + G2, trainData)
summary(bestSub2fit)

```

```

## Call:
## lm(formula = G3 ~ failures + G2, data = trainData)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -9.7929 -0.3001 -0.0398  0.8301  5.8735
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.30731   0.23818  -1.29 0.197378
## failures    -0.33362   0.09670  -3.45 0.000593 ***
## G2          1.04339   0.01976  52.82 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.649 on 727 degrees of freedom
## Multiple R-squared:  0.8254, Adjusted R-squared:  0.8249
## F-statistic: 1719 on 2 and 727 DF, p-value: < 2.2e-16
pred = predict(bestSub2fit, testData)
sqrt(mean((testData$G3 - pred)^2))

## [1] 1.441658
coef(regfit.full, 9)

## (Intercept) PstatusT Fjobservices travelttime failures famsupyes
## -0.80204721 -0.30128171 -0.22236350  0.17420908 -0.33391953  0.30884703
## paidyes      absences       G1           G2
## -0.38027783  0.02280147  0.09592451  0.97846994

bestSub9fit <- lm(G3 ~ Pstatus + Fjob + travelttime + failures + famsup + paid + absences + G1 + G2, data = trainData)
summary(bestSub9fit)

##
## Call:
## lm(formula = G3 ~ Pstatus + Fjob + travelttime + failures + famsup +
##     paid + absences + G1 + G2, data = trainData)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -9.5557 -0.4579  0.1305  0.8408  5.5619
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.473595   0.419284  -1.130 0.259051
## PstatusT     -0.318071   0.196979  -1.615 0.106806
## Fjobhealth   -0.440912   0.417212  -1.057 0.290956
## Fjobother    -0.393856   0.256652  -1.535 0.125325
## Fjobservices -0.612581   0.266094  -2.302 0.021614 *
## Fjobteacher   -0.717692   0.337035  -2.129 0.033559 *
## travelttime   0.180461   0.082118   2.198 0.028298 *
## failures     -0.341023   0.096196  -3.545 0.000418 ***
## famsupyes     0.318290   0.127282   2.501 0.012618 *
## paidyes      -0.377849   0.154392  -2.447 0.014630 *

```

```

## absences      0.022881   0.009506   2.407 0.016331 *
## G1           0.101446   0.038498   2.635 0.008593 **
## G2           0.978772   0.034441  28.419 < 2e-16 ***
##
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.619 on 717 degrees of freedom
## Multiple R-squared:  0.8341, Adjusted R-squared:  0.8313
## F-statistic: 300.4 on 12 and 717 DF,  p-value: < 2.2e-16
pred = predict(bestSub9fit, newdata = testData)
sqrt(mean((testData$G3 - pred)^2))

## [1] 1.433748

• Backward Selection

fullFit = lm(G3 ~ ., trainData)
backward_model = step(fullFit, direction = "backward", trace = F)
summary(backward_model)

##
## Call:
## lm(formula = G3 ~ Pstatus + travelttime + failures + famsup +
##     paid + absences + G1 + G2, data = trainData)
##
## Residuals:
##    Min      1Q  Median      3Q      Max
## -9.7161 -0.4484  0.1183  0.7976  5.4141
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.845936  0.358734 -2.358 0.01863 *
## PstatusT    -0.337606  0.196082 -1.722 0.08554 .
## travelttime  0.176224  0.081627  2.159 0.03119 *
## failures    -0.341178  0.096272 -3.544 0.00042 ***
## famsupyes   0.318997  0.126149  2.529 0.01166 *
## paidyes     -0.388332  0.154445 -2.514 0.01214 *
## absences     0.022916  0.009527  2.405 0.01640 *
## G1           0.098088  0.038365  2.557 0.01077 *
## G2           0.976470  0.034464  28.333 < 2e-16 ***
##
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.623 on 721 degrees of freedom
## Multiple R-squared:  0.8324, Adjusted R-squared:  0.8305
## F-statistic: 447.6 on 8 and 721 DF,  p-value: < 2.2e-16
preds = predict(backward_model, testData)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.419982

• Forward Selection

nullFit = lm(G3 ~ 1, trainData)
fullFit = lm(G3 ~ ., trainData)

```

```

forward_model = step(nullFit, scope = list(lower = nullFit, upper = fullFit), direction = "forward", trace = TRUE)
summary(forward_model)

##
## Call:
## lm(formula = G3 ~ G2 + failures + G1 + absences + paid + famsup +
##     travelttime + Pstatus, data = trainData)
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -9.7161 -0.4484  0.1183  0.7976  5.4141 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.845936  0.358734 -2.358  0.01863 *  
## G2          0.976470  0.034464 28.333 < 2e-16 *** 
## failures   -0.341178  0.096272 -3.544  0.00042 *** 
## G1          0.098088  0.038365  2.557  0.01077 *  
## absences    0.022916  0.009527  2.405  0.01640 *  
## paidyes    -0.388332  0.154445 -2.514  0.01214 *  
## famsupyes  0.318997  0.126149  2.529  0.01166 *  
## travelttime 0.176224  0.081627  2.159  0.03119 *  
## PstatusT   -0.337606  0.196082 -1.722  0.08554 .  
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
##
## Residual standard error: 1.623 on 721 degrees of freedom
## Multiple R-squared:  0.8324, Adjusted R-squared:  0.8305 
## F-statistic: 447.6 on 8 and 721 DF,  p-value: < 2.2e-16

preds = predict(forward_model, testData)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.419982

```

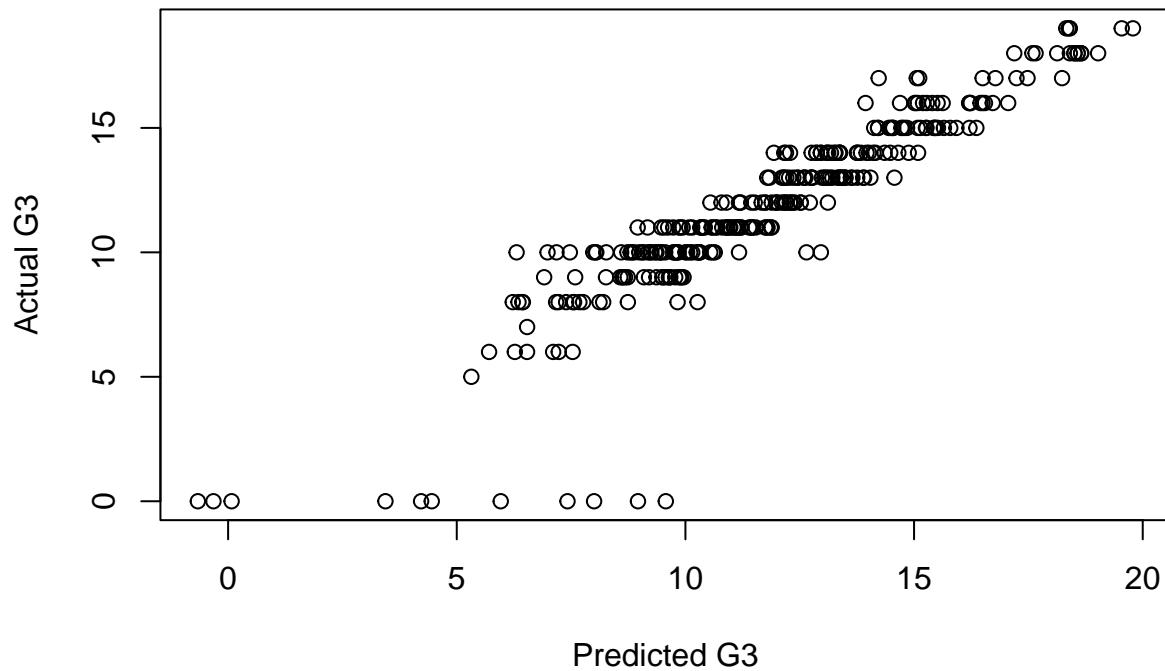
- The forward and backward selection give the same output so it would be the same rmse.

```

plot(preds, testData$G3, xlab = "Predicted G3", ylab = "Actual G3", main = "Predicted vs Actual Final G3")

```

Predicted vs Actual Final Grades



- Noticed that the number of absences increases the predicted final grade, which should not be the case, so we should a model without the absences.

```
forward_model_noAbsence = lm(formula = G3 ~ G2 + failures + G1 + paid + famsup +
    travelttime + Pstatus, data = trainData)
summary(forward_model_noAbsence)
```

```
##
## Call:
## lm(formula = G3 ~ G2 + failures + G1 + paid + famsup + travelttime +
##     Pstatus, data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -9.8243 -0.4128  0.1273  0.8095  5.3394 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.68230   0.35339 -1.931  0.053910 .  
## G2          0.97703   0.03458 28.256 < 2e-16 *** 
## failures    -0.32195   0.09626 -3.345  0.000866 *** 
## G1          0.09522   0.03847  2.475  0.013557 *  
## paidyes     -0.34518   0.15391 -2.243  0.025214 *  
## famsupyes   0.31719   0.12656  2.506  0.012425 *  
## travelttime  0.17145   0.08187  2.094  0.036602 *  
## PstatusT    -0.38252   0.19584 -1.953  0.051177 .  
## ---
```

```

## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.628 on 722 degrees of freedom
## Multiple R-squared:  0.831, Adjusted R-squared:  0.8294
## F-statistic: 507.3 on 7 and 722 DF,  p-value: < 2.2e-16
preds = predict(forward_model_noAbsence, testData)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.430864

• Try ridge regression

#install.packages("glmnet")
library(glmnet)

## Loading required package: Matrix

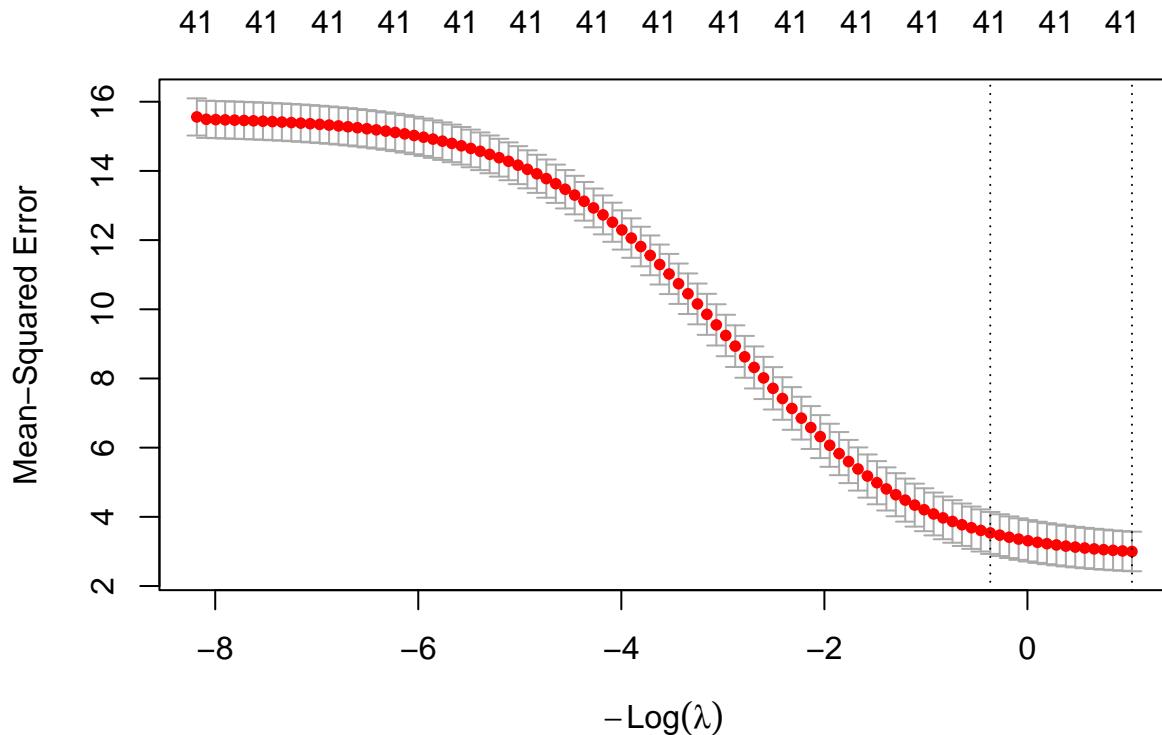
## Loaded glmnet 4.1-10

x = model.matrix(G3 ~ ., all_data) [, -1]
y = all_data$G3
grid = 10^seq(10, -2, length = 100)
ridge.mod = glmnet(x, y, alpha = 0, lambda = grid)
dim(coef(ridge.mod))

## [1] 42 100

set.seed(1)
cv.out = cv.glmnet(x[trainIdx, ], y[trainIdx], alpha = 0)
plot(cv.out)

```

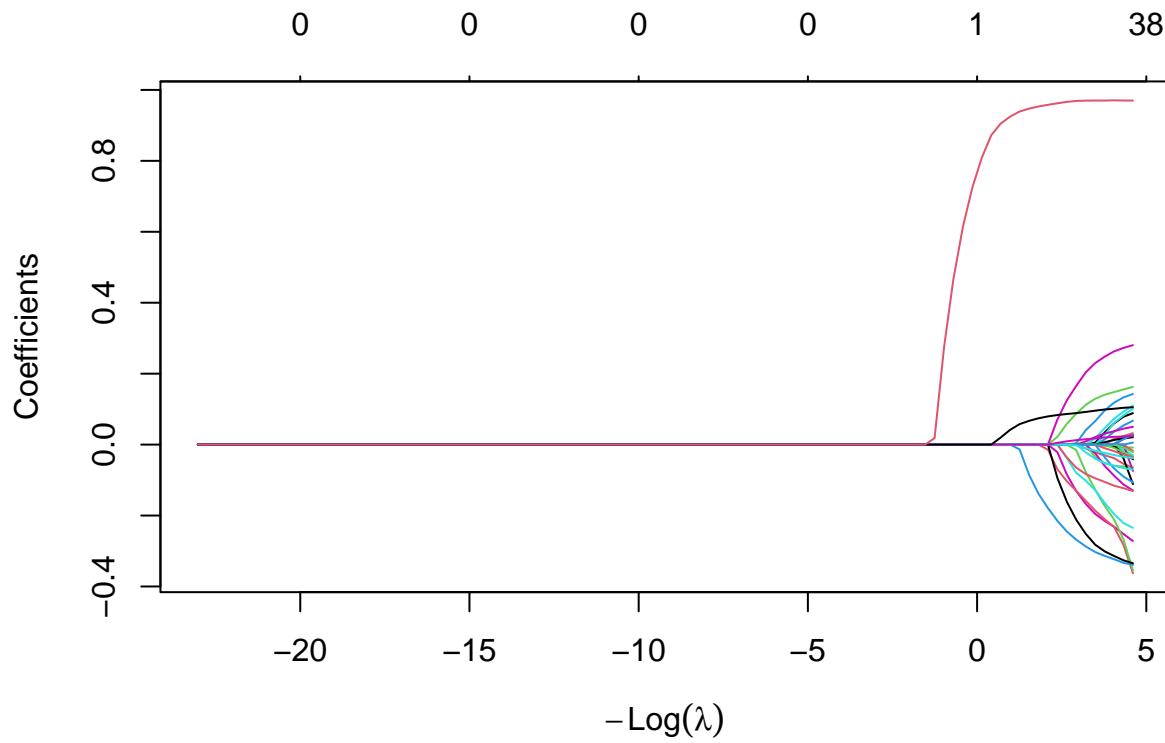


```

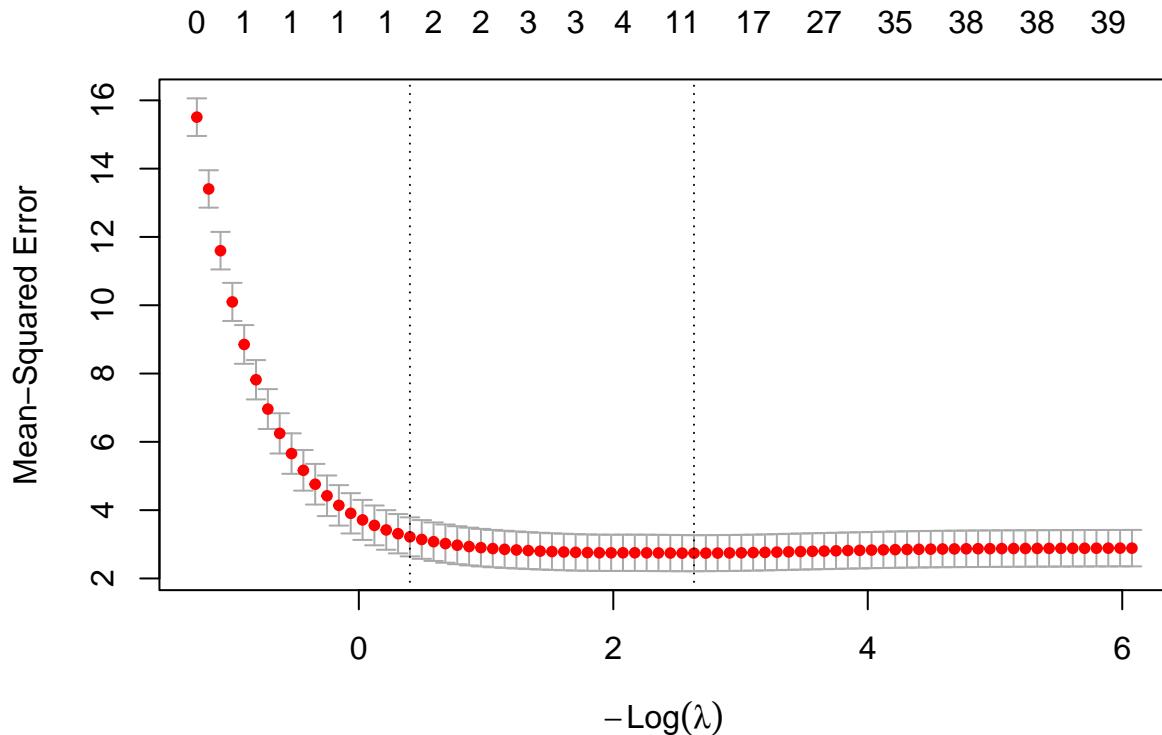
bestlam = cv.out$lambda.min
bestlam
## [1] 0.3572372
ridge.pred = predict(ridge.mod, s = bestlam, newx = x[-trainIdx, ])
sqrt(mean((ridge.pred - y[-trainIdx])^2))

## [1] 1.408071
lasso.mod = glmnet(x[trainIdx, ], y[trainIdx], alpha = 1, lambda = grid)
plot(lasso.mod)

```



```
set.seed(1)
cv.out = cv.glmnet(x[trainIdx, ], y[trainIdx], alpha = 1)
plot(cv.out)
```



```

bestlam = cv.out$lambda.min
bestlam
## [1] 0.07177728

lasso.pred = predict(lasso.mod, s = bestlam, newx = x[-trainIdx, ])
sqrt(mean((lasso.pred - y[-trainIdx])^2))

## [1] 1.413576

• lasso gives a close rsme to the best 5 subset but it is still worse then it. We will now try different models.

out = glmnet(x, y, alpha = 1, lambda = grid)
lasso.coef = predict(out, type = "coefficients", s = bestlam)[1:(nFeatures + 1), ]
lasso.coef

##      (Intercept)      schoolMS       sexM        age
## -0.60825163  0.00000000  0.00000000  0.00000000
##      addressU     famsizeLE3     PstatusT       Medu
##  0.00000000  0.00000000 -0.01148485  0.00000000
##      Fedu      Mjobhealth     Mjobother   Mjobservices
##  0.00000000  0.00000000  0.00000000  0.00000000
##      Mjobteacher    Fjobhealth     Fjobother   Fjobservices
##  0.00000000  0.00000000  0.00000000 -0.03553269
##      Fjobteacher    reasonhome reasonother reasonreputation
##  0.00000000  0.00000000  0.00000000  0.00000000
##      guardianmother    guardianother    traveltimes    studytime
##  0.00000000  0.00000000  0.02242678  0.00000000

```

```

##      failures    schoolsupyes    famsupyes    paidyes
## -0.17726388     0.00000000    0.01082511   -0.16640612
## activtiesyes    nurseryyes    higheryes    internetyes
## 0.00000000     0.00000000    0.00000000    0.00000000
## romanticyes    0.00000000
##         浪漫主义

lasso.coef[lasso.coef != 0]

## (Intercept)      PstatusT Fjobservices    traveltime    failures    famsupyes
## -0.60825163  -0.01148485  -0.03553269    0.02242678   -0.17726388   0.01082511
## paidyes
## -0.16640612

xTest = model.matrix(G3 ~ ., testData)[ , -1]
pred = predict(out, s = bestlam, newx = xTest)
sqrt(mean((testData$G3 - pred)^2))

## [1] 1.400898

-Decision Tree

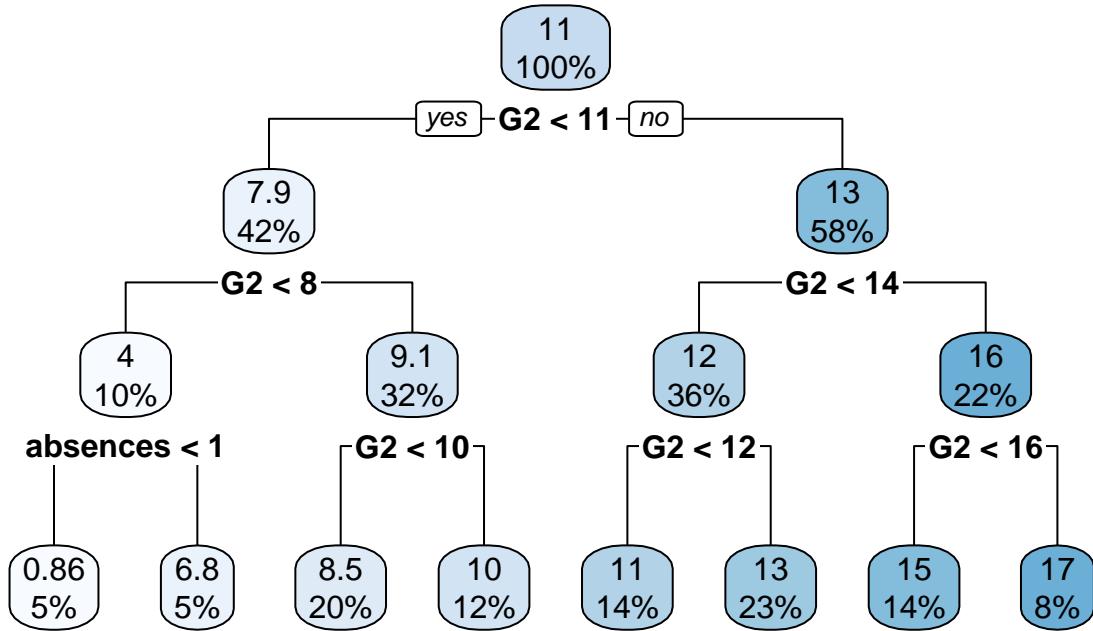
library(rpart)
library(rpart.plot)
tree_model = rpart(G3 ~ ., data = trainData, method = "anova")
print(tree_model)

## n= 730
##
## node), split, n, deviance, yval
##       * denotes terminal node
##
## 1) root 730 11325.87000 11.1178100
##    2) G2< 10.5 306 3781.67300 7.8562090
##      4) G2< 7.5 76 937.94740 3.9736840
##        8) absences< 1 36 214.30560 0.8611111 *
##        9) absences>=1 40 60.97500 6.7750000 *
##      5) G2>=7.5 230 1319.54800 9.1391300
##        10) G2< 9.5 146 894.49320 8.5068490 *
##        11) G2>=9.5 84 265.23810 10.2381000 *
##      3) G2>=10.5 424 1939.66000 13.4717000
##        6) G2< 13.5 264 384.35980 12.1325800
##          12) G2< 11.5 99 66.02020 11.1414100 *
##          13) G2>=11.5 165 162.72730 12.7272700 *
##        7) G2>=13.5 160 300.74370 15.6812500
##          14) G2< 15.5 101 73.80198 14.8910900 *
##          15) G2>=15.5 59 55.93220 17.0339000 *

rpart.plot(tree_model, main = "Decision Tree for Student Performance")

```

Decision Tree for Student Performance



```

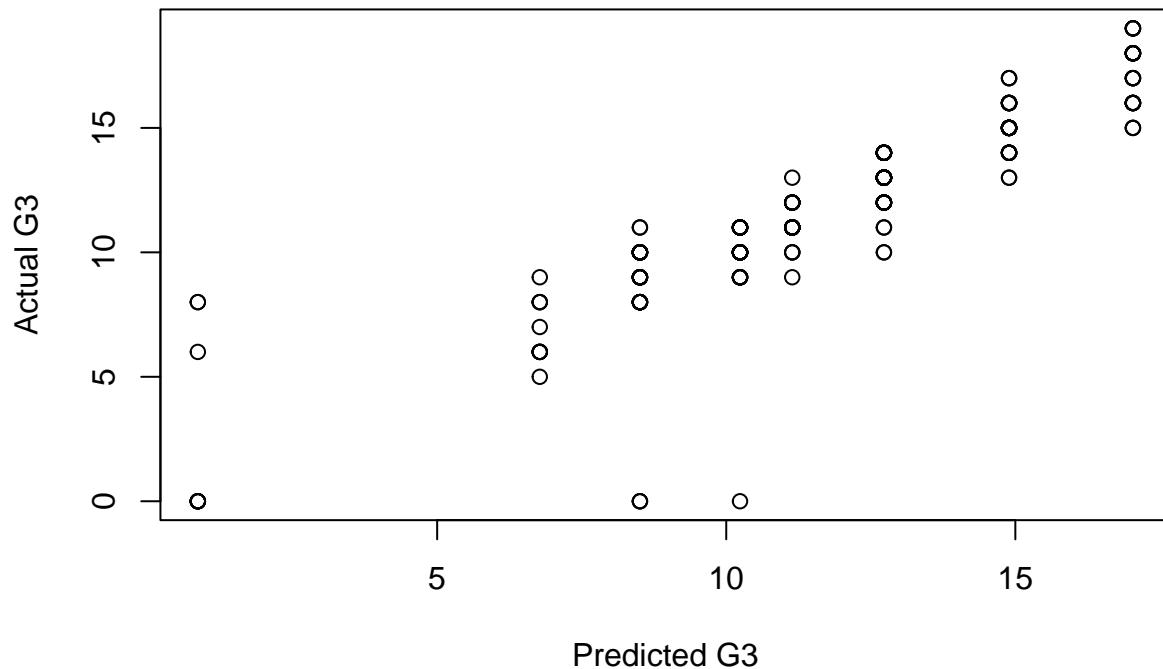
tree_pred = predict(tree_model, newdata = testData)
tree_rmse = sqrt(mean((testData$G3 - tree_pred)^2))
cat("Decision Tree RMSE:", tree_rmse, "\n")
  
```

Decision Tree RMSE: 1.532331

```

plot(tree_pred, testData$G3, xlab = "Predicted G3", ylab = "Actual G3", main = "Predicted vs Actual Final Grade")
  
```

Predicted vs Actual Final Grades (Decision Tree)

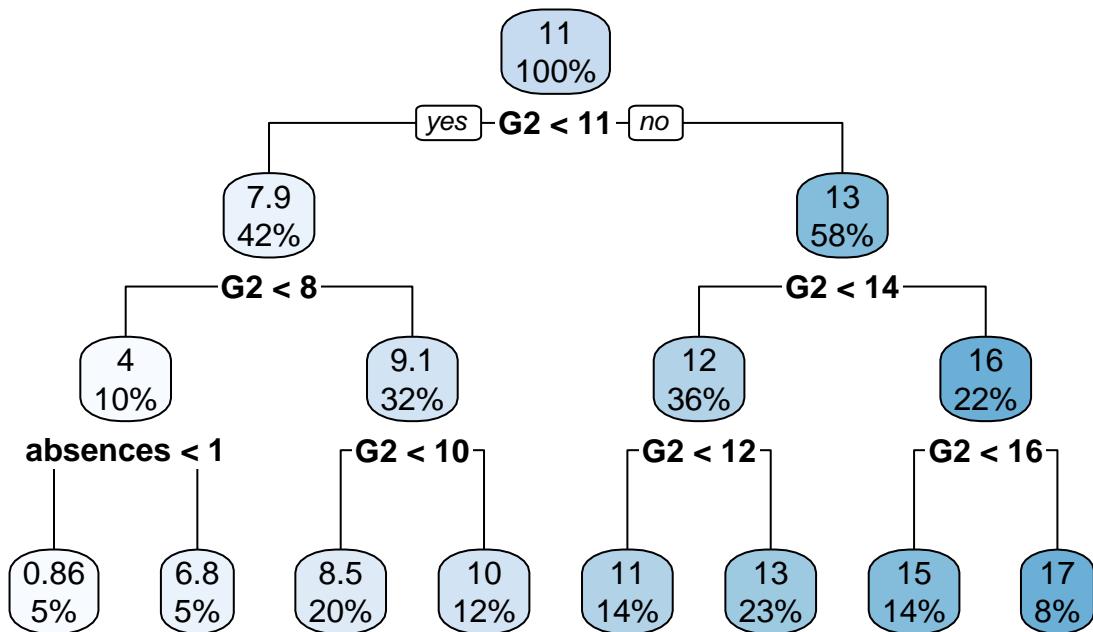


```

## rpart(formula = G3 ~ ., data = trainData, method = "anova")
##
## Variables actually used in tree construction:
## [1] absences G2
##
## Root node error: 11326/730 = 15.515
##
## n= 730
##
##          CP nsplit rel error  xerror      xstd
## 1 0.494844      0    1.00000 1.00197 0.069863
## 2 0.134575      1    0.50516 0.51508 0.036031
## 3 0.110769      2    0.37058 0.39734 0.028796
## 4 0.058509      3    0.25981 0.26677 0.026359
## 5 0.015099      4    0.20130 0.20851 0.026481
## 6 0.014111      5    0.18620 0.19409 0.026306
## 7 0.013740      6    0.17209 0.18688 0.026532
## 8 0.010000      7    0.15835 0.16998 0.025799
optimal_cp = tree_model$cptable[which.min(tree_model$cptable[, "xerror"]), "CP"]
pruned_tree = prune(tree_model, cp = optimal_cp)
rpart.plot(pruned_tree, main = "Pruned Decision Tree for Student Performance")

```

Pruned Decision Tree for Student Performance



```

pruned_pred = predict(pruned_tree, newdata = testData)
pruned_rmse = sqrt(mean((testData$G3 - pruned_pred)^2))
cat("Pruned Decision Tree RMSE:", pruned_rmse, "\n")

```

```

## Pruned Decision Tree RMSE: 1.532331

Lets try using the two the two features that the decision tree used for a linear regression It was worse then
the best

twoFeaturelm = lm(G3 ~ absences + G2, data = trainData)

twoFeaturePreds = predict(twoFeaturelm, testData)
sqrt(mean((testData$G3 - twoFeaturePreds)^2))

## [1] 1.424654

#install.packages("randomForest")
library(randomForest)

## randomForest 4.7-1.2

## Type rfNews() to see new features/changes/bug fixes.

set.seed(1)
bag.grade = randomForest(G3 ~ ., trainData, mtry = nFeatures, importance = TRUE)
bag.grade

##
## Call:
##   randomForest(formula = G3 ~ ., data = trainData, mtry = nFeatures,      importance = TRUE)
##   Type of random forest: regression
##   Number of trees: 500
##   No. of variables tried at each split: 32
##
##   Mean of squared residuals: 2.578824
##   % Var explained: 83.38

summary(bag.grade)

##          Length Class  Mode
## call           5   -none- call
## type          1   -none- character
## predicted     730  -none- numeric
## mse           500  -none- numeric
## rsq           500  -none- numeric
## oob.times     730  -none- numeric
## importance    64   -none- numeric
## importanceSD  32   -none- numeric
## localImportance 0   -none- NULL
## proximity     0   -none- NULL
## ntree          1   -none- numeric
## mtry           1   -none- numeric
## forest         11  -none- list
## coefs          0   -none- NULL
## y              730  -none- numeric
## test           0   -none- NULL
## inbag          0   -none- NULL
## terms          3   terms  call

bag.preds = predict(bag.grade, testData)
sqrt(mean((testData$G3 - bag.preds)^2))

## [1] 1.410638

```

```
nFeatures
```

```
## [1] 32
```

Do CV to find the best number of variables to consider

```
k = 5
```

```
set.seed(1)
```

```
folds = sample(rep(1:k, length = nTrain))
```

```
cv.errors = matrix(NA, k, nFeatures, dimnames = list(NULL, paste(1:nFeatures)))
```

```
for (j in 1:k) {
```

```
  train_fold = trainData[folds != j, ]
```

```
  val_fold = trainData[folds == j, ]
```

```
  for (i in 1:nFeatures) {
```

```
    cvModel = randomForest(G3 ~ ., data = train_fold, mtry = i, importance = TRUE)
```

```
    preds = predict(cvModel, val_fold)
```

```
    cv.errors[j, i] = sqrt(mean((val_fold$G3 - preds)^2))
```

```
}
```

```
}
```

```
mean.cv.errors = apply(cv.errors, 2, mean)
```

```
mean.cv.errors
```

```
##      1       2       3       4       5       6       7       8
```

```
## 3.031255 2.500724 2.218642 2.036103 1.926500 1.841903 1.784860 1.736578
```

```
##      9      10      11      12      13      14      15      16
```

```
## 1.699038 1.670485 1.656379 1.646655 1.636822 1.627610 1.621393 1.612126
```

```
##     17      18      19      20      21      22      23      24
```

```
## 1.616449 1.596153 1.611678 1.601829 1.601387 1.603941 1.598686 1.608574
```

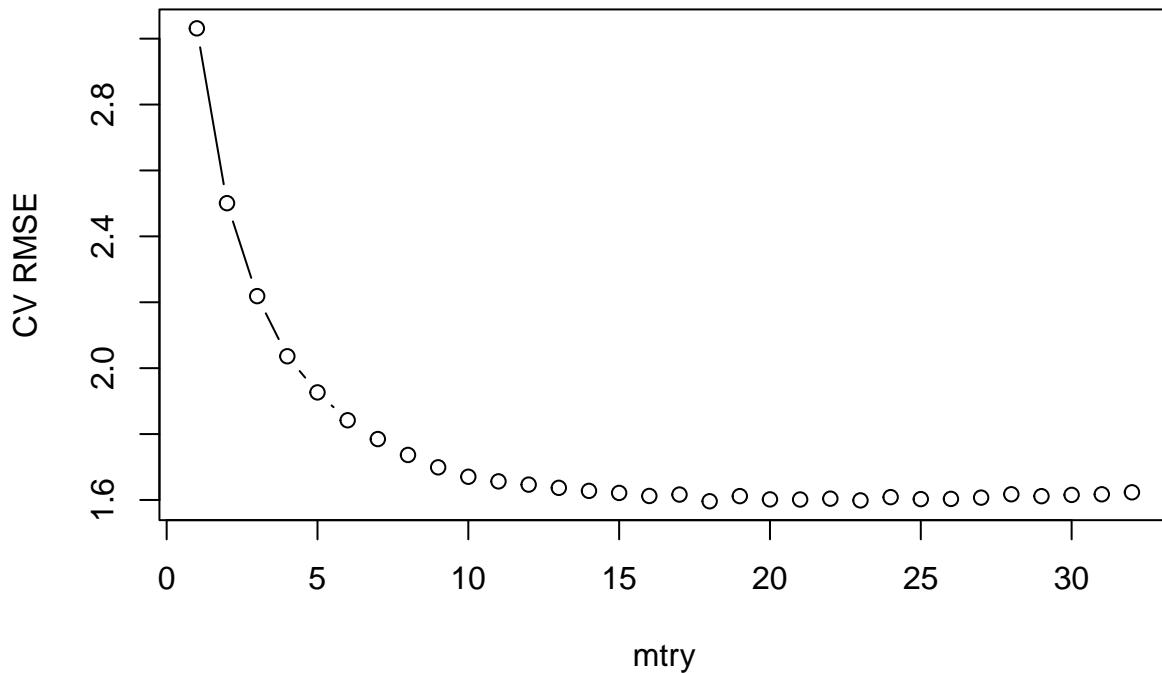
```
##     25      26      27      28      29      30      31      32
```

```
## 1.602603 1.603243 1.607095 1.617302 1.611540 1.615390 1.617281 1.623264
```

```
par(mfrow = c(1,1))
```

```
plot(1:nFeatures, mean.cv.errors, type = "b", xlab = "mtry", ylab = "CV RMSE", main = "5-Fold CV for Ra
```

5-Fold CV for Random Forest mtry



```
best_mtry = which.min(mean.cv.errors)
best_mtry

## [1] 18
## [1] 18

best_mtry_model = randomForest(G3 ~ ., data = trainData, mtry = best_mtry, importance = TRUE)
preds = predict(best_mtry_model, testData)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.376701

5-fold cv for ntree
k = 5
set.seed(1)

num_tree = c(300,400,500,600,700,800,900,1000,1100,1200,1300,1400,1500)
folds = sample(rep(1:k, length = nTrain))
cv.errors = matrix(NA, k, length(num_tree), dimnames = list(NULL, paste(num_tree)))

for (j in 1:k) {
  train_fold = trainData[folds != j, ]
  val_fold = trainData[folds == j, ]
  for (i in 1:length(num_tree)) {
    cvModel = randomForest(G3 ~ ., data = train_fold, mtry = best_mtry, ntree = num_tree[i], importance
    preds = predict(cvModel, val_fold)
    cv.errors[j, i] = sqrt(mean((val_fold$G3 - preds)^2))
```

```

    }
}

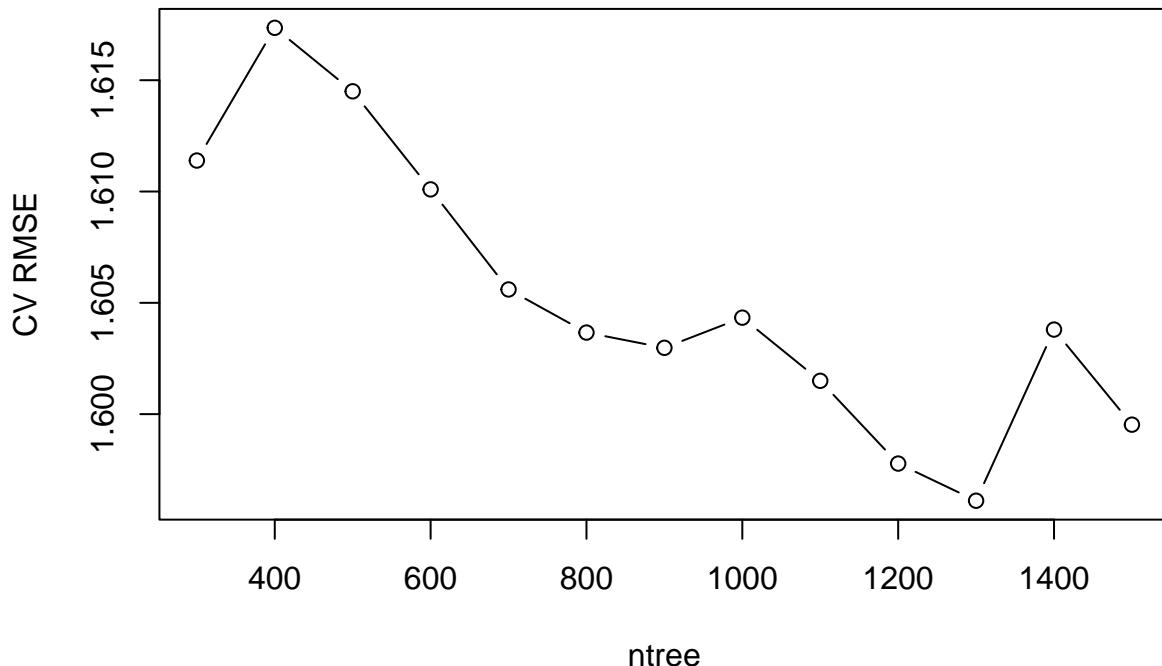
mean.cv.errors = apply(cv.errors, 2, mean)
mean.cv.errors

##      300      400      500      600      700      800      900      1000
## 1.611392 1.617354 1.614502 1.610096 1.605600 1.603665 1.602978 1.604334
##      1100     1200     1300     1400     1500
## 1.601495 1.597779 1.596110 1.603798 1.599526

set.seed(1)
best_mtry = 18
par(mfrow = c(1,1))
plot(num_tree, mean.cv.errors, type = "b", xlab = "ntree", ylab = "CV RMSE", main = "5-Fold CV for Random Forest with mtry = 18")

```

5–Fold CV for Random Forest ntree with mtry = 18



```

best_ntree = num_tree[which.min(mean.cv.errors)]
best_ntree

## [1] 1300

best_mtry_ntree_model = randomForest(G3 ~ ., data = trainData, mtry = best_mtry, ntree = best_ntree, importance = TRUE)
preds = predict(best_mtry_ntree_model, testData)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.3666702

```

```

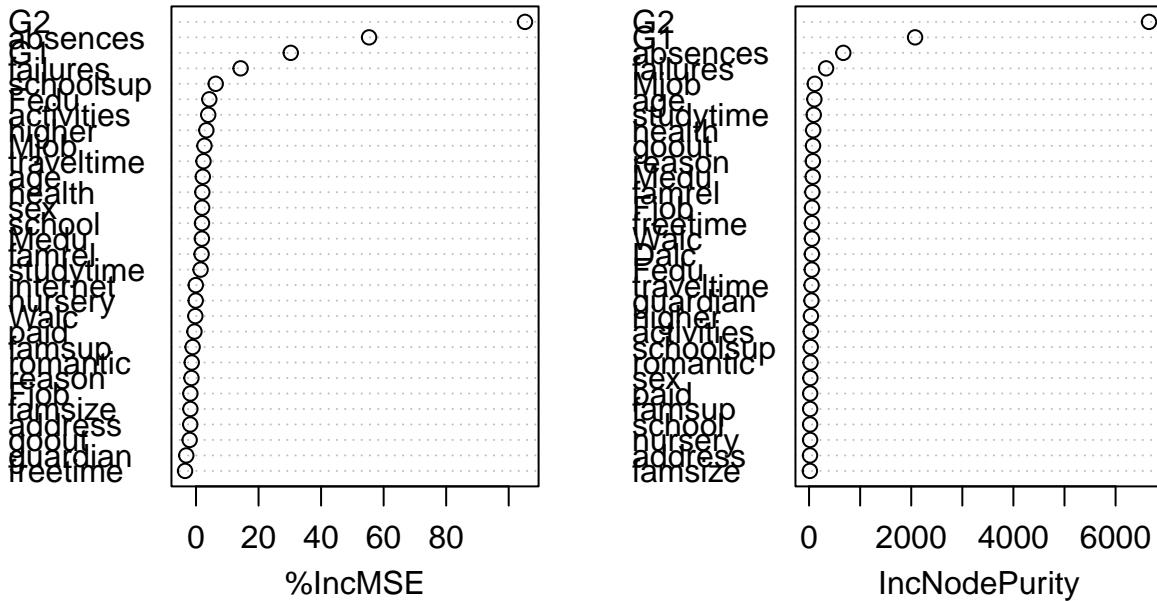
importance(best_mtry_ntree_model)

## %IncMSE IncNodePurity
## school      1.92094201    20.541576
## sex         1.96421389    25.294556
## age          2.17470514   104.078247
## address     -1.83767197   18.487838
## famsize     -1.80755593   15.973303
## Pstatus     -3.76015357    6.642354
## Medu        1.87883590   71.921914
## Fedu        4.25690818   52.998818
## Mjob         2.70042893  112.209146
## Fjob        -1.79803929   57.381035
## reason      -1.45861900   74.740303
## guardian    -3.11023809   46.798700
## traveltim   2.39247331   49.250670
## studytime   1.46361460   92.476565
## failures    14.27141506  333.475364
## schoolsup   6.30013103   32.399816
## famsup      -1.08890145   21.256512
## paid         -0.49984802   22.049340
## activities  3.90476080   33.147660
## nursery     -0.09296337   20.518667
## higher       3.30714037   36.405934
## internet    -0.07230990   14.718405
## romantic    -1.36424888   29.196371
## famrel       1.78780560   64.745881
## freetime    -3.50070337   56.861962
## goout        -2.02397436   77.551666
## Dalc         -5.03100753   54.966561
## Walc        -0.22194110   56.149429
## health       2.02594611   80.793628
## absences     55.34964853  669.920630
## G1           30.29733386  2075.478004
## G2           105.21044518  6654.193589

```

```
varImpPlot(best_mtry_ntree_model)
```

best_mtry_ntree_model



try only using the best 4 features that are used G2, G1, absences and failures and the current best ntree

```
set.seed(1)
mostImportFeaturesRandomForeset = c('G2', 'G1', 'absences', 'failures', 'G3')
mostImportTrain = trainData[, mostImportFeaturesRandomForeset]
mostImportTest = testData[, mostImportFeaturesRandomForeset]
mostImportModel = randomForest(G3 ~ ., mostImportTrain, mtry = ncol(mostImportTrain) - 1, ntrees = best)
preds = predict(mostImportModel, mostImportTest)
sqrt(mean((mostImportTest$G3 - preds)^2))
```

```
## [1] 1.544647
```

Try k-fold cv for number of trees using this subset when only consider these features for splits

```
k = 5
set.seed(1)

num_tree = c(300, 400, 500, 600, 700, 800, 900, 1000, 1100, 1200, 1300, 1400, 1500)
folds = sample(rep(1:k, length = nTrain))
cv.errors = matrix(NA, k, length(num_tree), dimnames = list(NULL, paste(num_tree)))

for (j in 1:k) {
  train_fold = mostImportTrain[folds != j, ]
  val_fold = mostImportTrain[folds == j, ]
  for (i in 1:length(num_tree)) {
    cvModel = randomForest(G3 ~ ., data = train_fold, mtry = 4, ntree = num_tree[i], importance = TRUE)
    preds = predict(cvModel, val_fold)
    cv.errors[j, i] = sqrt(mean((val_fold$G3 - preds)^2))
```

```

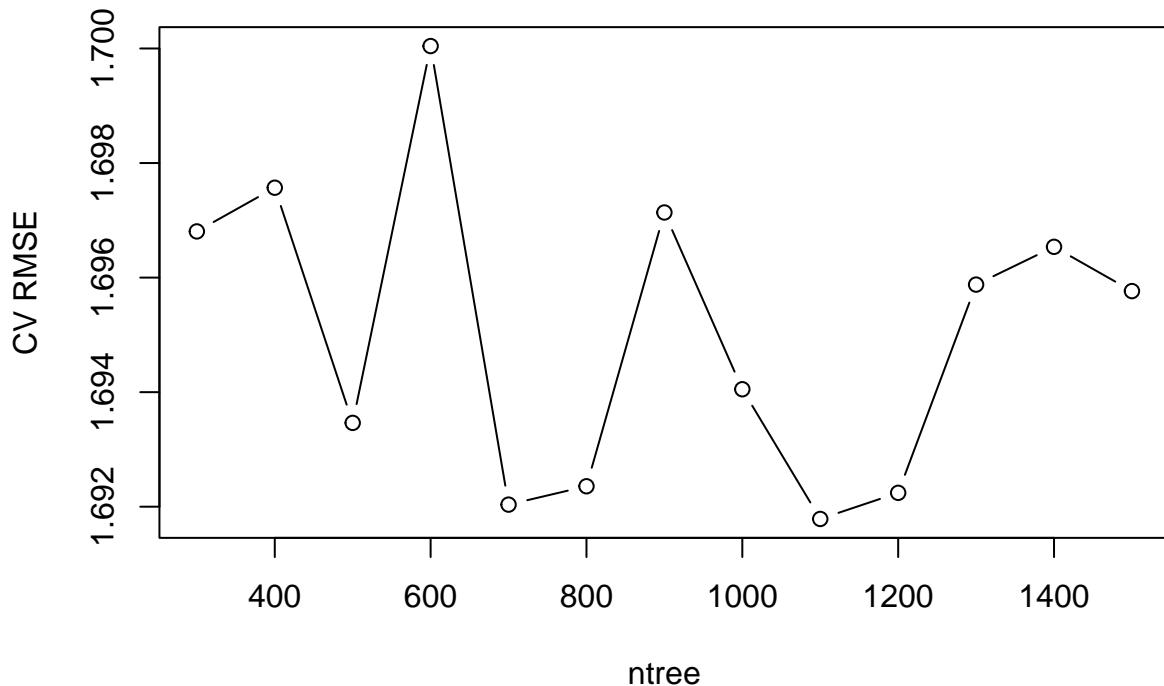
    }
}

mean.cv.errors = apply(cv.errors, 2, mean)
mean.cv.errors

##      300      400      500      600      700      800      900      1000
## 1.696806 1.697569 1.693463 1.700042 1.692037 1.692357 1.697137 1.694051
##      1100     1200     1300     1400     1500
## 1.691786 1.692242 1.695878 1.696537 1.695764
par(mfrow = c(1,1))
plot(num_tree, mean.cv.errors, type = "b", xlab = "ntree", ylab = "CV RMSE", main = "5-Fold CV for Random Forest ntrees")

```

5-Fold CV for Random Forest ntrees with Important Features



```

best_ntree = num_tree[which.min(mean.cv.errors)]
best_ntree

## [1] 1100

best_mtry_ntree_model = randomForest(G3 ~ ., data = mostImportTrain, mtry = 4, ntree = best_ntree, importance = TRUE)
preds = predict(best_mtry_ntree_model, mostImportTest)
sqrt(mean((testData$G3 - preds)^2))

## [1] 1.526788

```